Middleware and Web Services

Lecture 2: Application Protocols

doc. Ing. Tomáš Vitvar, Ph.D.
tomas@vitvar.com • @TomasVitvar • http://vitvar.com

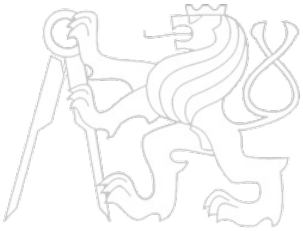

Czech Technical University in Prague
Faculty of Information Technologies • Software and Web Engineering • http://vitvar.com/courses/mdw

Modified: Tue Jan 30 2018, 10:26:51
Humla v0.3

Overview

- Introduction to Application Protocols
  - *Synchronous and Asynchronous Communication*
  - *Selected Networking Concepts*
- Introduction to HTTP

Application Protocols

- Remember this

Drawing not available

- App protocols mostly on top of the TCP Layer
  - *use TCP socket for communication*
- Major protocols
  - *HTTP – most of the app protocols layered on HTTP*
    - → *wide spread, but: implementors often break HTTP semantics*
  - *RMI – Remote Method Invocation*
    - → *Java-specific, rather interface*
    - → *may use HTTP underneath (among other things)*
  - *XML-RPC – Remote Procedure Call and SOAP*
    - → *Again, HTTP underneath*
  - *WebSocket – new protocol part of HTML5*

Socket

Drawing not available

- Handshaking (connection establishment)
  - *The server listens at* `[dst_ip,dsp_port]`
  - *Three-way handshake:*
    - → *the client at* `[src_ip,src_port]` *sends a connection request*
    - → *the server responds*
    - → *the client acknowledges the response, can send data along*
  - *Result is a socket (virtual communication channel) with unique identification:* `socket=[src_ip,src_port;dst_ip,dst_port]`
- Data transfer (resource usage)
  - *Client/server writes/reads data to/from the socket*
  - *TCP features: reliable delivery, correct order of packets, flow control*
- Connection close

Addressing in Application Protocol

Drawing not available

- IP addressing: IP is an address of a machine interface
  - *A machine can have multiple interfaces (eth0, eth1, bond0, ...)*
- TCP addressing: TCP port is an address of an app running on a machine and listening on a machine interface
  - *Multiple applications with different TCP ports may listen on a machine interface*
- Application addressing
  - *Additional mechanisms to address entities within an application*
  - *They are out of scope of IP/TCP, they are app specific*
    - → *for example, Web apps served by a single Web server*

Virtual IP

Drawing not available

- Virtual IP
  - *Additional IP addresses assigned to a network interface*
    - → *For example,* `eth0 - eth0:1, eth0:2, eth0:3, ...`
    - → *A process can bind to the virtual IP*
    - → *Multiple processes can listen on the same tcp port but on different virtual IPs*
- Benefits
  - *Floating IP – a process can move transparently to another physical machine*
  - *Network configuration can be preserved, no need to reconfigure*
  - *Failover concept uses floating IPs*

Virtual IP Configuration

- Steps to configure virtual IP in Linux (example for `eth0`)
  1. *Find out the interface's network mask*

```
$ ifconfig eth0
eth0      Link encap:Ethernet  HWaddr 00:0C:29:AB:5E:6A
inet addr:172.16.169.184  Bcast:172.16.169.255  Mask:255.255.255.0
...
```

  2. *Create virtual IP using* `ifconfig`
     - *it should use the same network mask*
     - *it should be free, usually allocated to be used as a virtual IP*

```
$ sudo ifconfig eth0:1 172.16.169.186 netmask 255.255.255.0
$ ifconfig eth0:1
eth0:1    Link encap:Ethernet  HWaddr 00:0C:29:AB:5E:6A
          inet addr:172.16.169.186  Bcast:172.16.169.255  Mask:255.255.25
```

  3. *Update neighbours' ARP (Address Resolution Protocol) caches*
     - *to associate the virtual IP with MAC address of* `eth0`
     - *when the virtual IP was in use on other node or interface*

```
$ sudo arping -q -U -c 3 -I eht0 172.16.169.184
```

Tasks

- *Configure a virtual IP on your computer and test it using* `ping`

Server Failover

- Failover
  - *Failover = ability to relocate the server to another machine without an impact on the performance*

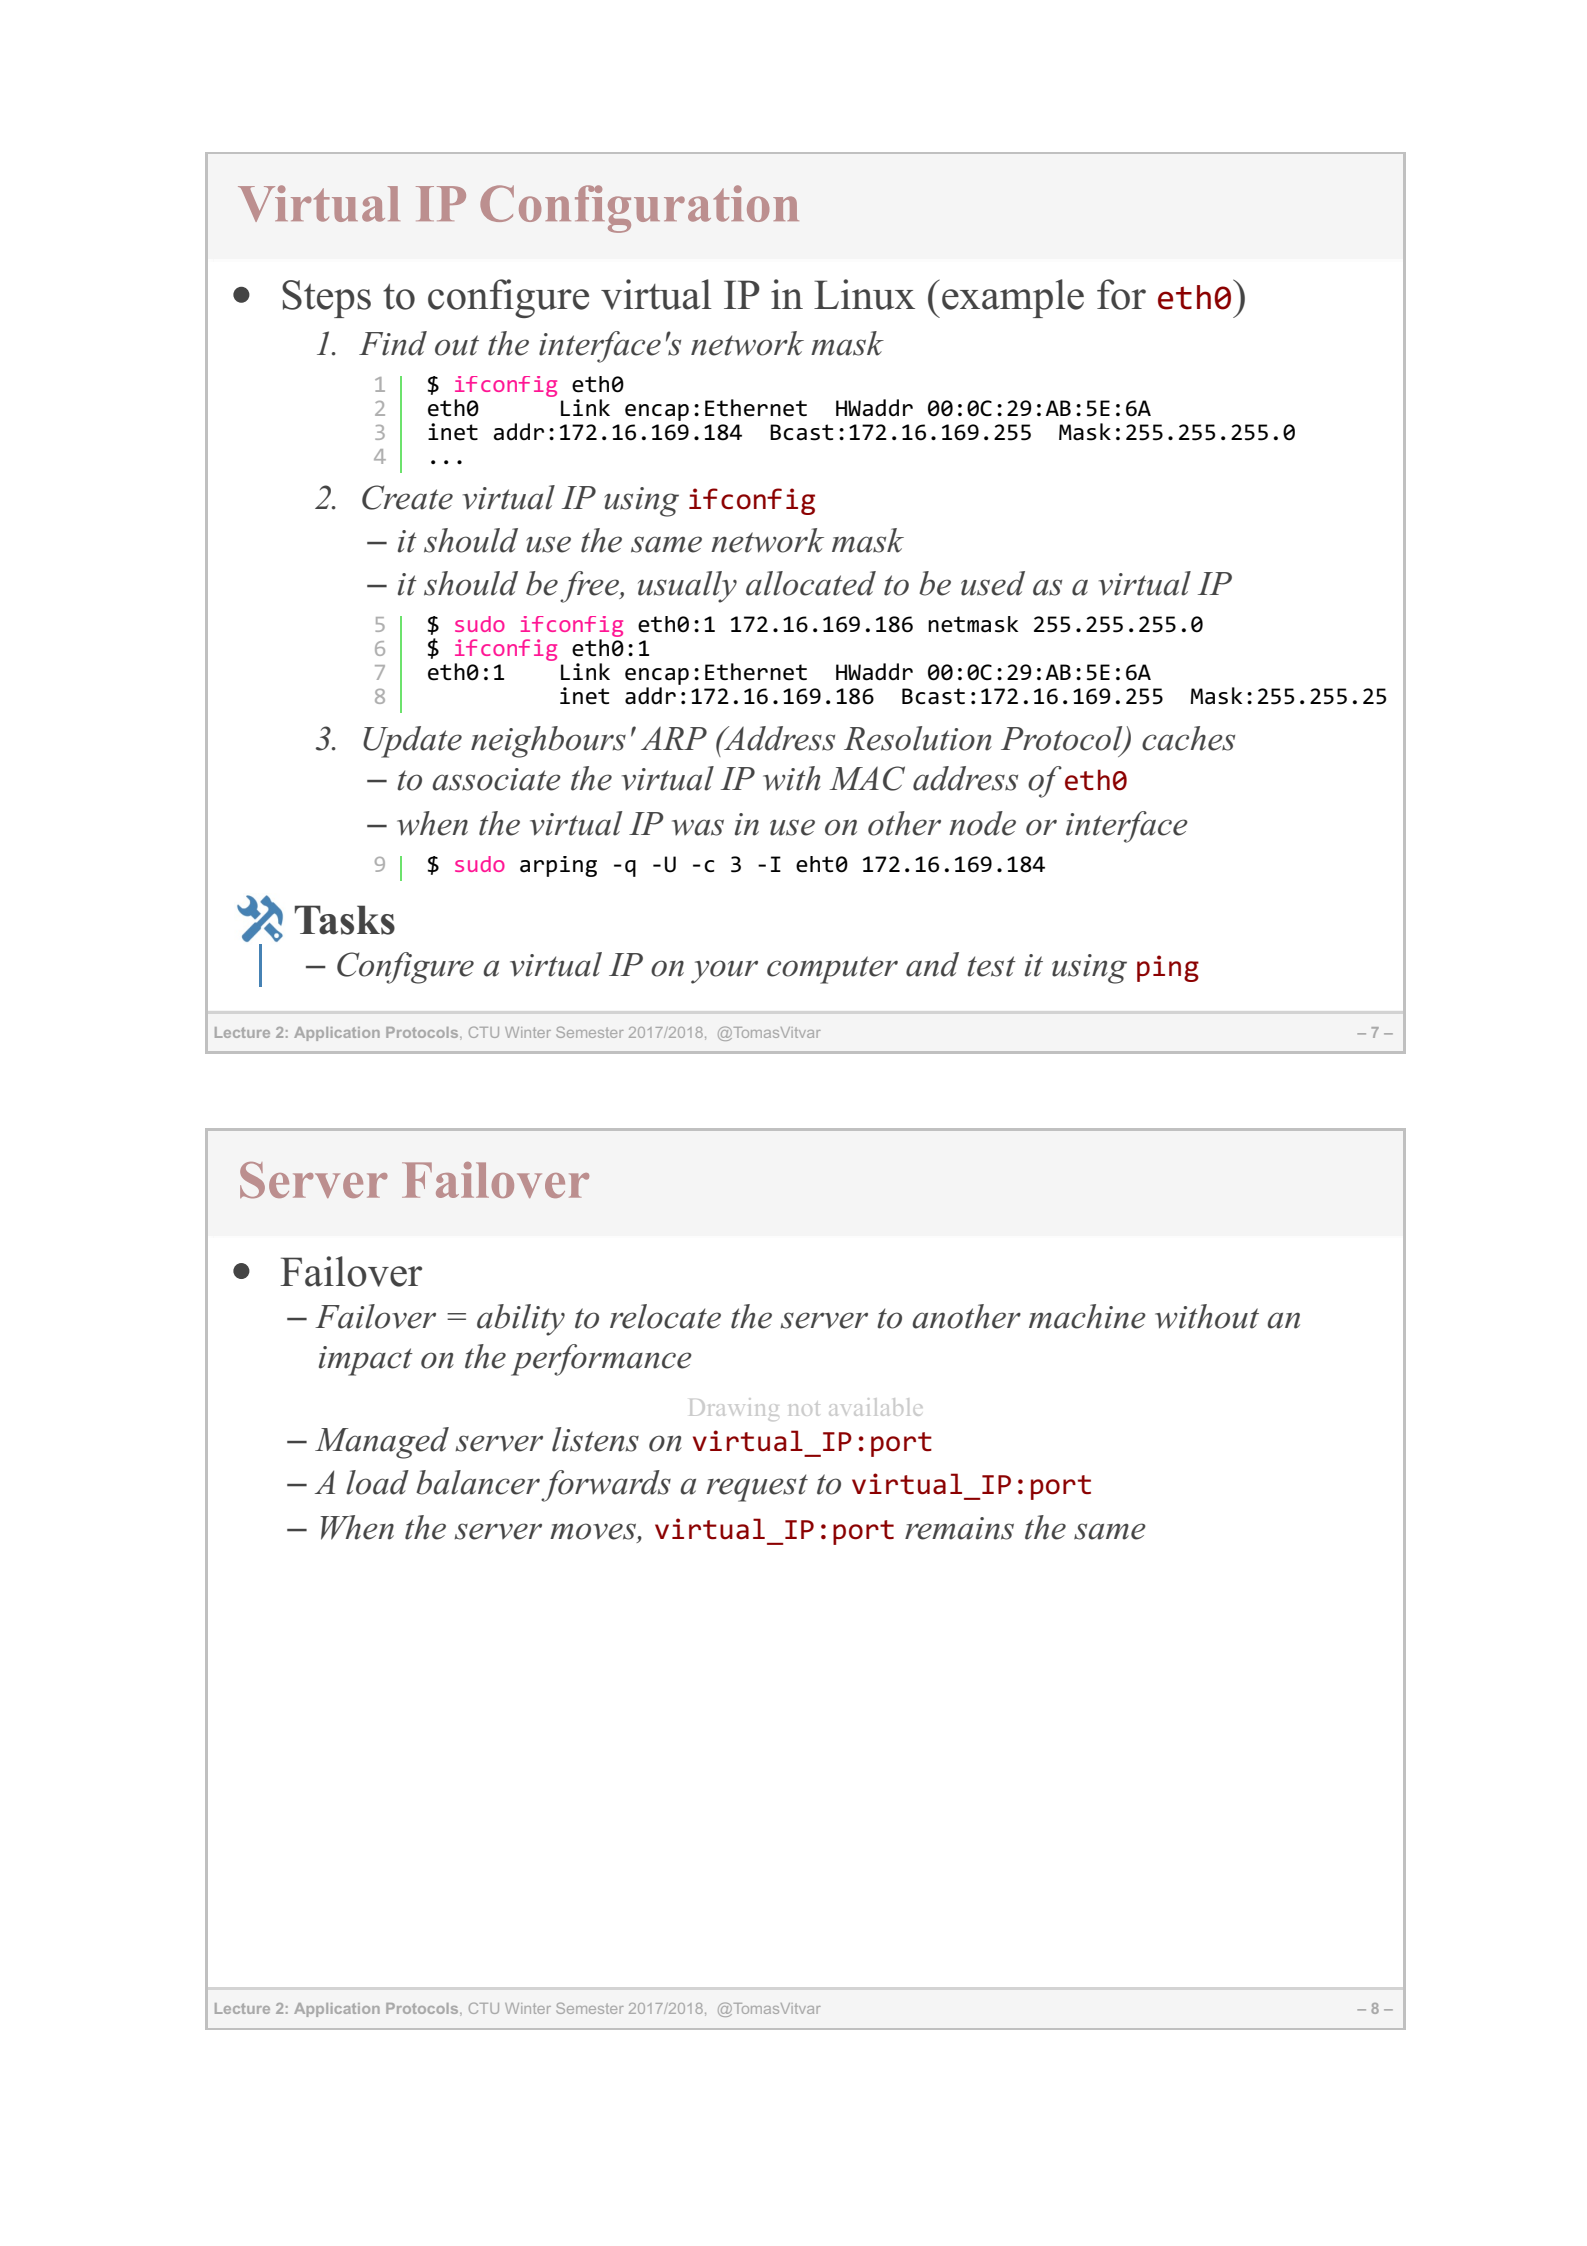

  - *Managed server listens on* `virtual_IP:port`
  - *A load balancer forwards a request to* `virtual_IP:port`
  - *When the server moves,* `virtual_IP:port` *remains the same*

Overview

- Introduction to Application Protocols
  - *Synchronous and Asynchronous Communication*
  - *Selected Networking Concepts*
- Introduction to HTTP

Synchronous and Asynchronous Communication

Drawing not available

- Synchronous
  - *one socket, $|t_{req} - t_{res}|$ is small*
  - *easy to implement and deploy, only standard firewall config*
  - *only the server defines endpoint*
- Asynchronous
  - *request, response each has socket, client and server define endpoints*
  - *$|t_{req} - t_{res}|$ can be large (hours, even days)*
  - *harder to do across network elements (private/public networks issue)*

Asynchronous via Intermediary

Drawing not available

- Intermediary
  - *A component that decouples a client-server communication*
  - *It increases reliability and performance*
    - → *The server may not be available when a client sends a request*
    - → *There can be multiple servers that can handle the request*
- Further Concepts
  - *Message Queues (MQ) – queue-based communication*
  - *Publish/Subscribe (P/S) – event-driven communication*

Asynchronous via Polling

Drawing not available

- Polling – only clients open sockets
  - *A client performs multiple request-response interactions*
    - → *The first interaction initiates a process on the server*
    - → *Subsequent interactions check for the processing status*
    - → *The last interaction retrieves the processing result*
- Properties of environments
  - *A server cannot open a socket with the client (network restrictions)*
  - *Typically on the Web (a client runs in a browser)*

Overview

- Introduction to Application Protocols
  - *Synchronous and Asynchronous Communication*
  - *Selected Networking Concepts*
- Introduction to HTTP

Public/Private Network Configuration

Drawing not available

- Adds complexity to configuration of application
  - *Config example at server with* `eth0 = 147.32.100.1` *(iptables)*

```
1  # enable ip forwarding from one interface to another within linux core
2  echo 1 > /proc/sys/net/ipv4/ip_forward
3
4  # redirect all communication coming to tcp/3000 to 192.168.1.2:4000
5  iptables -t nat -A PREROUTING -i eth0 -p tcp --dport 3000 -j DNAT \
6      --to-dest 192.168.1.2 --to-port 4000
```

Virtual Private Network

Drawing not available

- VPN = Virtual Private Network
  - *an overlay network between a client and a server*
  - *the network spans accross underlying network elements*
  - *Example:*
    - → *VPN client starts a VPN connection with the VPN server via network interfaces*
    - → *VPN server assigns an IP address to the VPN client from the server's subnet*
    - → *Packets in VPN communication are encrypted and sent out in an outer VPN*

Forward Proxy

Drawing not available

- Forward Proxy
  - *Centralized access control based on content*
  - *The client knows about the site it wants to access*
  - *Perfoms request on behalf of the client*
    - → *Caches content to increase performance, limits network traffic*
    - → *Filters requests or controls access based on destinations or origins*
  - *Widely used in private networks in companies*
  - *Most of the proxy servers today are Web proxy servers*

Reverse Proxy

Drawing not available

- Reverse Proxy
  - *Aggregates multiple request-response interactions with back-end systems*
  - *Processes the request on behalf of the client*
  - *The client does not know about the back-end systems*
  - *May provide additional capabilities*
    - → *Data transformations*
    - → *Security – authentication, authorization*
    - → *Orchestration of communication with back-end systems*

Overview

- Introduction to Application Protocols
- Introduction to HTTP
  - *State Management*

Hypertext Transfer Protocol – HTTP

- Application protocol, basis of Web architecture
  - *Part of HTTP, URI, and HTML family*
  - *Request-response protocol*
- One socket for single request-response
  - *original specification*
  - *have changed due to performance issues*
    - → *many concurrent requests*
    - → *overhead when establishing same connections*
    - → *HTTP 1.1 offers persistent connection and pipelining*
- HTTP is stateless
  - *Multiple HTTP requests cannot be normally related at the server*
    - → *"problems" with state management*
    - → *REST goes back to the original HTTP idea*

HTTP Request and Response

- Request Syntax

```
method uri http-version <crlf>
(header : value <crlf>)*
<crlf>
[ data ]
```

- Response Syntax

```
http-version response-code [ message ]  <crlf>
(header : value <crlf>)*
<crlf>
[ data ]
```

- Semantics of terms

```
method         = "GET" | "POST" | "DELETE" | "PUT" | "HEAD" | "OPTIONS"
uri            = [ path ] [ ";" params ] [ "?" query ]
http-version   = "HTTP/1.0" | "HTTP/1.1"
response-code  = valid response code
header : value = valid HTTP header and its value
data           = resource state representation (hypertext)
```

Serving HTTP Request

Drawing not available

- IP and TCP addressing
  1. *User enters URL* `http://company.cz:8080/orders` *to the browser*
  2. *Browser gets an IP address for* `company.cz`, `IP:138.232.189.127`
  3. *Browser and Web Server creates a socket* `[147.32.100.5:3223;138.232.189.127:8080]`
- Application addressing
  4. *Browser sends HTTP request, that is, writes following data to the socket*

```
1 GET /orders HTTP/1.1
2 Host: company.cz
```

  5. *Web server passes the request to the web application* `company.cz` *which serves* `GET orders` *and that writes a response back to the socket.*

HTTP Listener

- HTTP listener implementation in Java using Jetty
  - *Server listens on port* `8080`
  - *Jetty parses HTTP request data into* `HttpServletRequest` *object.*
  - *When a client connects, the method* `handleRequest` *is called*
  - *The method tests the value of the* `host` *header and responds back if the header matches* `company.cz` *value.*

```
1  /** handles the request when client connects **/
2  public void handleRequest(HttpServletRequest request,
3          HttpServletResponse response) throws IOException, ServletException {
4
5      // test if the host is company.cz
6      if (request.getHeader("Host").equals("company.cz")) {
7          response.setStatus(200);
8          response.setHeader("Content-Type", "text/plain");
9          response.getWriter().write("This is the response");
10         response.flushBuffer();
11     } else
12         response.sendError(400); // bad request
13 }
```

HTTP Listener (Cont.)

- Test it using Telnet

```
telnet 127.0.0.1 8080
# ...lines omitted due to brevity
GET /orders HTTP/1.1
Host: company.cz

HTTP/1.1 201 OK
Content-Type: plain/text

This is the response...
```

- HTTP listener in bash
  - *Use it to test incomming HTTP connections quickly*
  - *Uses* nc *utility (netcat)*

```
# ctrl-c to stop http listener
control_c() {
  echo -en "\n* Exiting\n"
  exit $?
}
trap control_c SIGINT

for (( ; ; ))
do
        echo -e "\n\n* Listening on port $1..."
        echo -e "\nHTTP/1.0 204 No Content\n\n" | nc -l $port
done
```


Virtual Web Server

- Virtual server
  - *Configuration of a named virtual web server*
  - *Web server uses host request header to distinguish among multiple virtual web servers on a single physical host.*
- Apache virtual Web server configuration
  - *Two virtual servers hosted on a single physical host*

```
# all IP addresses will be used for named virtual hosts
NameVirtualHost *:80

<VirtualHost *:80>
        ServerName company.com
        ServerAdmin admin@company.com
        DocumentRoot /var/www/apache/company.com
</VirtualHost>

<VirtualHost *:80>
        ServerName firm.cz
        ServerAdmin admin@firm.cz
        DocumentRoot /var/www/apache/firm.cz
</VirtualHost>
```

Better Support for HTTP Testing

- Use `curl` to test HTTP protocol

```
Usage: curl [options...] <url>

-X/--request <command>          Specify request command to use
-H/--header <line>              Custom header to pass to server
-d/--data <data>                HTTP POST data
-b/--cookie <name=string/file>  Cookie string or file to read cookies from
-v/--verbose                    Make the operation more talkative
```

- Example

```
curl -v -H "Host: company.cz"  127.0.0.1:8080

* About to connect() to 127.0.0.1 port 8080
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 port 8080
> GET / HTTP/1.1
> User-Agent: curl/7.20.0 (i386-apple-darwin10.3.2) libcurl/7.20.0 OpenSSL/0.9
> Accept: */*
> Host: company.cz
>
< HTTP/1.1 201 OK
< Connection: keep-alive
< Content-Type: plain/text
<
< This is the response...
```


Overview

- Introduction to Application Protocols
- Introduction to HTTP
  - *State Management*

State Management

- HTTP is a stateless protocol – original design
  - *No information to relate multiple interactions at server-side*
    - → *Except* `Authorization` *header is copied in every request*
    - → *IP addresses do not work, one public IP can be shared by multiple clients*
- Solutions to check for a valid state at server-side
  - ***Cookies*** *– obvious and the most common workaround*
    - → *RFC 2109 – HTTP State Management Mechanism*
    - → *Allow clients and servers to talk in a context called* ***sessions***
  - ***Hypertext*** *– original HTTP design principle*
    - → *App states represented by resources (hypermedia), links define transitions between states*
    - → *Adopted by the REST principle* ***statelessness***

Interaction with Cookies

- Request-response interaction with cookies
  - *Session is a logical channel maintained by the server*

Drawing not available

- Stateful Server
  - *Server remembers the session information in a server memory*
  - *Server memory is a non-persistent storage, when server restarts the memory content is lost!*

Set-Cookie and Cookie Headers

- `Set-Cookie` response header

```
set-cookie = "Set-Cookie:" cookie ("," cookie)*
  cookie    = NAME "=" VALUE (";" cookie-av)*
  cookie-av = "Comment" "=" value
            | "Domain" "=" value
            | "Max-Age" "=" value
            | "Path" "=" value
```

  - `domain` – *a domain for which the cookie is applied*
  - `Max-Age` – *number of seconds the cookie is valid*
  - `Path` – *URL path for which the cookie is applied*

- `Cookie` request header. A client sends the cookie in a request if:
  - `domain` *matches the origin server's fully-qualified host name*
  - `path` *matches a prefix of the request-URI*
  - `Max-Age` *has not expired*

```
cookie  = "Cookie:" cookie-value (";" cookie-value)*
  cookie-value    = NAME "=" VALUE [";" path] [";" domain]
  path            = "$Path" "=" value
  domain          = "$Domain" "=" value
```

  - `domain`, *and* `path` *are values from corresponding attributes of the* `Set-Cookie` *header*

Session Management Java Class

- Manages client sessions in a server memory

```
public class Sessions<E> {

    // storage for the session data;
    private Hashtable<String, E> sessions = new Hashtable<String, E>();

    /** Returns session id based on the information in the http request **/
    public String getSessionID(HttpServletRequest request) throws Exception {
        String sid = null;

        // extract the session id from the cookie
        if (request.getHeader("cookie") != null) {
            Pattern p = Pattern.compile(".*session-id=([a-zA-Z0-9]+).*");
            Matcher m = p.matcher(request.getHeader("cookie"));
            if (m.matches()) sid = m.group(1);
        }

        // create the session id md5 hash; use random number to generate a client-id
        // note that this is a simple solution but not very reliable
        if (sid == null || sessions.get(sid) == null) {
            MessageDigest md = MessageDigest.getInstance("MD5");
            md.update(new String(request.getRemoteAddr() +
                Math.floor(Math.random()*1000)).getBytes());
            sid = Utils.toHexString(md.digest());
        }
        return sid;
    }

    public E getData(String sid) ... // returns session data from sessions object
    public void setData(String sid, E d) ... // sets session data to sessions object
}
```

Stateful Server Implementation

- Simple per-client counter

```java
public void handleRequest(HttpServletRequest request,
        HttpServletResponse response) throws Exception {
    // get the session id
    String sid = sessions.getSessionID(request);

    // create the new data if none exists
    if (sessions.getData(sid) != null)
        sessions.setData(sid,
            Integer.valueOf(sessions.getData(sid).intValue() + 1));
    else
        sessions.setData(sid, Integer.valueOf(1));

    // send the response
    response.setStatus(200);
    response.setHeader("Set-Cookie", "session-id="+ sid + "; MaxAge=3600");
    response.setHeader("Content-Type", "text/plain");
    response.getWriter().write("Number of hits from you: " +
        sessions.getData(sid).toString());
    response.flushBuffer();
}
```

Task

– *What happens when the server restarts?*

– *How do you change the code to count requests from all clients?*

Testing

- Testing
  – `curl` *will require you to specify cookies in every request*
  – *Browser handles cookies automatically*

```
# run curl for the first time
curl -v 127.0.0.1:8080
> GET / HTTP/1.1
> Host: 127.0.0.1:8080
>
< HTTP/1.1 200 OK
< Set-Cookie: session-id=3a9c3cdc5ff36434aa1ba860727ca401;max-age=3600
<
Number of hits from you: 1

# copy the cookie session-id from previous response
curl -v -b session-id=3a9c3cdc5ff36434aa1ba860727ca401 127.0.0.1:8080
> GET / HTTP/1.1
> Host: 127.0.0.1:9900
> Cookie: session-id=3a9c3cdc5ff36434aa1ba860727ca401
>
< HTTP/1.1 200 OK
< Set-Cookie: session-id=3a9c3cdc5ff36434aa1ba860727ca401;max-age=3600
<
Number of hits from you: 2
```