

Middleware Architectures 1

Lecture 3: Communication Protocols

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <http://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <http://vitvar.com/courses/mdw>



Modified: Sat Sep 19 2020, 17:17:08
Humla v0.3

Overview

- **Introduction to Application Protocols**
 - *Synchronous and Asynchronous Communication*
 - *Introduction to HTTP*

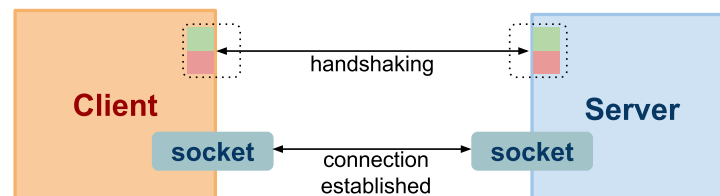
Application Protocols

- Remember this



- App protocols mostly on top of the TCP Layer
 - use TCP socket for communication
- Major protocols
 - HTTP – most of the app protocols layered on HTTP
 - widely spread
 - RMI – Remote Method Invocation
 - Java-specific; vendor-interoperability problem
 - may use HTTP underneath (among other things)
 - XML-RPC and SOAP – Remote Procedure Call and SOAP
 - HTTP-based
 - WebSocket – new protocol part of HTML5

Socket



- Handshaking (connection establishment)
 - The server listens at `[dst_ip, dst_port]`
 - Three-way handshake:
 - the client sends a connection request with TCP flags (SYN, $x=rand$)
 - the server responds with its own TCP flags (SYN ACK, $x+1$ $y=rand$)
 - the client acknowledges the response, can send data along (ACK, $y+1$ $x+1$)
 - Result is a socket (virtual communication channel) with unique identification:


```
socket=[src_ip,src_port;dst_ip,dst_port]
```
- Data transfer (resource usage)
 - Client/server writes/reads data to/from the socket
 - TCP features: reliable delivery, correct order of packets, flow control
- Connection close

New Connection Costs

- Creating a new TCP connection is expensive
 - It requires to complete a full roundtrip
 - It is limited by a network latency, not bandwidth
- Example
 - Distance from London to New York is approx. 5500 km
 - Communication over a fibre link will take at least 28ms one way
 - Three-way handshake will take a minimum of 56ms
- Connection reuse is critical for any app running over TCP
 - HTTP Keep-alive
 - HTTP pipelining
- TCP Fast Open (TFO)
 - TFO allows to speed up the opening of successive TCP connections
 - TCP cookie stored on the client that was established on initial connection
 - The client sends the TCP cookie with SYN packet
 - The server verifies the TCP cookie and can send the data without final ACK
 - Can reduce network transaction latency by 15%
 - TFO is supported by Linux in 3.7+ kernels

Addressing in Application Protocol

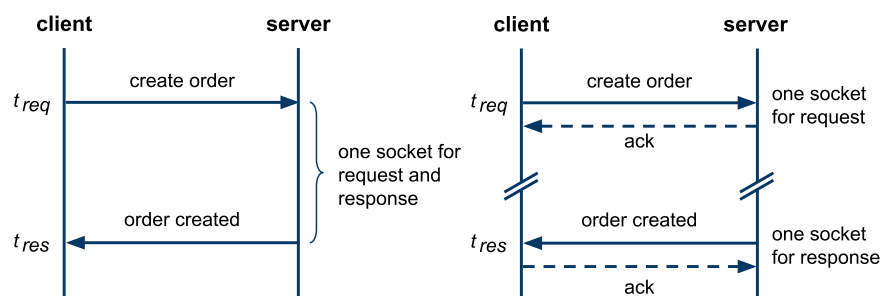


- IP addressing: IP is an address of a machine interface
 - A machine can have multiple interfaces (eth0, eth1, bond0, ...)
- TCP addressing: TCP port is an address of an app running on a machine and listening on a machine interface
 - Multiple applications with different TCP ports may listen on a machine interface
- Application addressing
 - Additional mechanisms to address entities within an application
 - They are out of scope of IP/TCP, they are app specific
 - for example, Web apps served by a single Web server

Overview

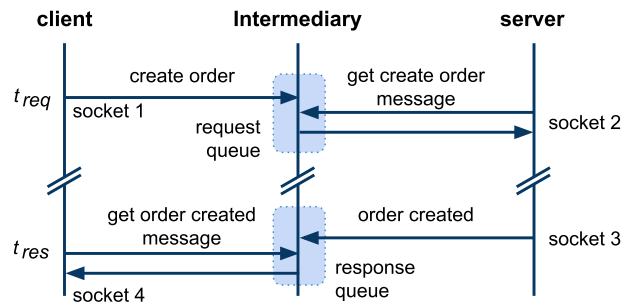
- Introduction to Application Protocols
 - *Synchronous and Asynchronous Communication*
 - *Introduction to HTTP*

Synchronous and Asynchronous Communication



- Synchronous
 - one socket, $|t_{req} - t_{res}|$ is small
 - easy to implement and deploy, only standard firewall config
 - only the server defines endpoint
- Asynchronous
 - request, response each has socket, client and server define endpoints
 - $|t_{req} - t_{res}|$ can be large (hours, even days)
 - harder to do across network elements (private/public networks issue)

Asynchronous via Intermediary



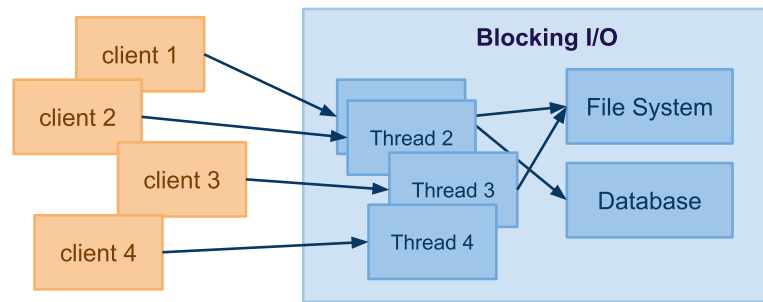
- **Intermediary**
 - A component that decouples a client-server communication
 - It increases reliability and performance
 - The server may not be available when a client sends a request
 - There can be multiple servers that can handle the request
- **Further Concepts**
 - Message Queues (MQ) – queue-based communication
 - Publish/Subscribe (P/S) – event-driven communication

Asynchronous via Polling



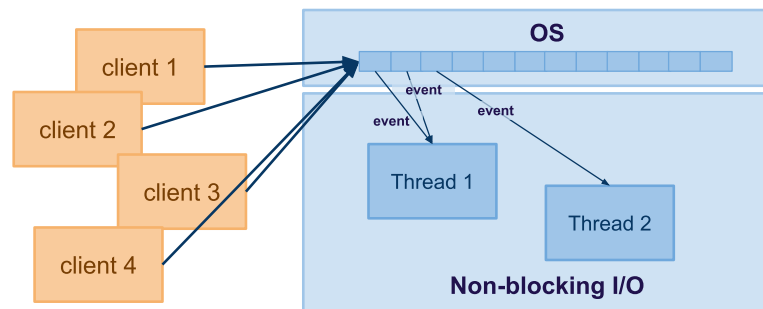
- **Polling** – only clients open sockets
 - A client performs multiple request-response interactions
 - The first interaction initiates a process on the server
 - Subsequent interactions check for the processing status
 - The last interaction retrieves the processing result
- **Properties of environments**
 - A server cannot open a socket with the client (network restrictions)
 - Typically on the Web (a client runs in a browser)

Blocking (Synchronous) I/O



- Inbound connection
 - A server creates a thread for every inbound connection
 - For example, 1K connections = 1K threads, big overhead
 - A thread is reserved for the entire duration of the request processing
- Outbound connection
 - A thread is blocked when outbound connection is made
 - When outbound connection is slow, the scalability is poor

Non-Blocking (Asynchronous) I/O



- Inbound connections
 - The connection is maintained by the OS, not the server app
 - The Web app registers events, OS triggers events when they occur
 - The app may create working threads and controls their number
- Outbound connections
 - The app registers a callback that is called when the data is available
 - Event loop

Overview

- Introduction to Application Protocols
 - *Synchronous and Asynchronous Communication*
 - *Introduction to HTTP*

Hypertext Transfer Protocol – HTTP

- Application protocol, basis of Web architecture
 - *Part of HTTP, URI, and HTML family*
 - *Request-response protocol*
- One socket for single request-response
 - *original specification*
 - *have changed due to performance issues*
 - *many concurrent requests*
 - *overhead when establishing same connections*
 - *HTTP 1.1 offers persistent connection and pipelining*
- HTTP is stateless
 - *Multiple HTTP requests cannot be normally related at the server*
 - *"problems" with state management*
 - *REST goes back to the original HTTP idea*

HTTP Request and Response

- Request Syntax

```
method uri http-version <crLf>
(header : value <crLf>)*
<crLf>
[ data ]
```

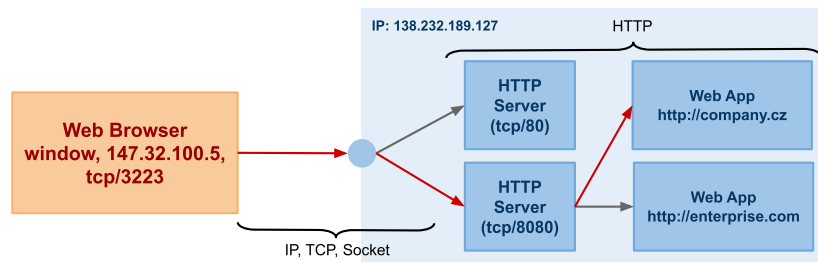
- Response Syntax

```
http-version response-code [ message ] <crLf>
(header : value <crLf>)*
<crLf>
[ data ]
```

- Semantics of terms

```
method          = "GET" | "POST" | "DELETE" | "PUT" | "HEAD" | "OPTIONS"
uri             = [ path ] [ ";" params ] [ "?" query ]
http-version    = "HTTP/1.0" | "HTTP/1.1"
response-code   = valid response code
header : value  = valid HTTP header and its value
data           = resource state representation (hypertext)
```

Serving HTTP Request



- IP and TCP addressing

1. User enters URL **http://company.cz:8080/orders** to the browser
2. Browser gets an IP address for **company.cz**, **IP:138.232.189.127**
3. Browser and Web Server creates a socket
[147.32.100.5:3223;138.232.189.127:8080]

- Application addressing

4. Browser sends HTTP request, that is, writes following data to the socket
 - 1 | GET /orders HTTP/1.1
 - 2 | Host: company.cz
5. Web server passes the request to the web application **company.cz** which serves **GET orders** and that writes a response back to the socket.

Virtual Web Server

- Virtual server
 - Configuration of a named virtual web server
 - Web server uses host request header to distinguish among multiple virtual web servers on a single physical host.
- Apache virtual Web server configuration
 - Two virtual servers hosted on a single physical host

```
1 # all IP addresses will be used for named virtual hosts
2 NameVirtualHost *:80
3
4 <VirtualHost *:80>
5     ServerName company.com
6     ServerAdmin admin@company.com
7     DocumentRoot /var/www/apache/company.com
8 </VirtualHost>
9
10 <VirtualHost *:80>
11     ServerName firm.cz
12     ServerAdmin admin@firm.cz
13     DocumentRoot /var/www/apache/firm.cz
14 </VirtualHost>
```

Better Support for HTTP Testing

- Use **curl** to test HTTP protocol

```
1 Usage: curl [options...] <url>
2
3 -X/--request <command>      Specify request command to use
4 -H/--header <line>          Custom header to pass to server
5 -d/--data <data>            HTTP POST data
6 -b/--cookie <name=string/file> Cookie string or file to read cookies from
7 -v/--verbose                Make the operation more talkative
```

- Example

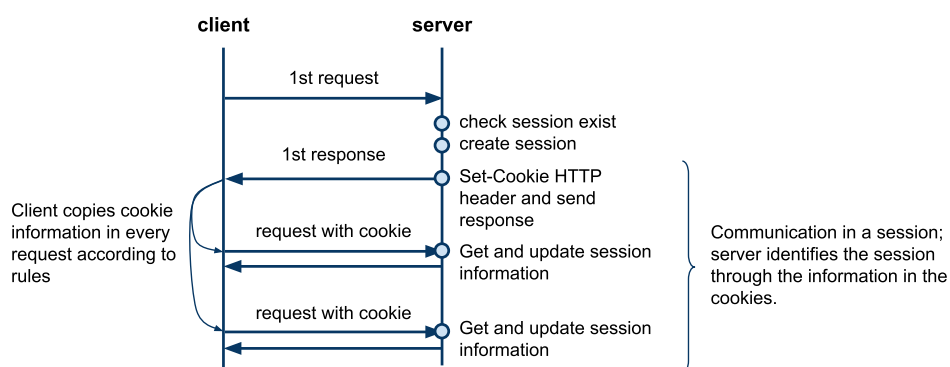
```
1 curl -v -H "Host: company.cz" 127.0.0.1:8080
2
3 * About to connect() to 127.0.0.1 port 8080
4 * Trying 127.0.0.1... connected
5 * Connected to 127.0.0.1 port 8080
6 > GET / HTTP/1.1
7 > User-Agent: curl/7.20.0 (i386-apple-darwin10.3.2) libcurl/7.20.0 OpenSSL/0.9.8
8 > Accept: */*
9 > Host: company.cz
10 >
11 < HTTP/1.1 201 OK
12 < Connection: keep-alive
13 < Content-Type: plain/text
14 <
15 < This is the response...
```

State Management

- HTTP is a stateless protocol – original design
 - No information to relate multiple interactions at server-side
 - Except **Authorization** header is copied in every request
 - IP addresses do not work, one public IP can be shared by multiple clients
- Solutions to check for a valid state at server-side
 - **Cookies** – obvious and the most common workaround
 - RFC 2109 – HTTP State Management Mechanism [🔗](#)
 - Allow clients and servers to talk in a context called **sessions**
 - **Hypertext** – original HTTP design principle
 - App states represented by resources (hypermedia), links define transitions between states
 - Adopted by the REST principle **statelessness**

Interaction with Cookies

- Request-response interaction with cookies
 - Session is a logical channel maintained by the server



- Stateful Server
 - Server remembers the session information in a server memory
 - Server memory is a non-persistent storage, when server restarts the memory content is lost!

Set-Cookie and Cookie Headers

- **Set-Cookie** response header

```
1 set-cookie = "Set-Cookie:" cookie ("," cookie)*
2 cookie     = NAME "=" VALUE (";" cookie-av)*
3 cookie-av  = "Comment" "=" value
4           | "Domain"   "=" value
5           | "Max-Age"  "=" value
6           | "Path"     "=" value
```

- **domain** – a domain for which the cookie is applied
- **Max-Age** – number of seconds the cookie is valid
- **Path** – URL path for which the cookie is applied

- **Cookie** request header. A client sends the cookie in a request if:

- **domain** matches the origin server's fully-qualified host name
- **path** matches a prefix of the request-URI
- **Max-Age** has not expired

```
1 cookie = "Cookie:" cookie-value (";" cookie-value)*
2 cookie-value = NAME "=" VALUE [";" path] [";" domain]
3 path        = "$Path" "=" value
4 domain      = "$Domain" "=" value
```

- **domain**, and **path** are values from corresponding attributes of the **Set-Cookie** header