

# Middleware Architectures 1

## Lecture 5: Cloud Native and Kubernetes

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <https://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <https://vitvar.com/lectures>



Modified: Sun Nov 09 2025, 21:18:03  
Humla v1.0

# Overview

- Cloud Native
- Kubernetes

# Overview

- The Cloud Native Computing Foundation (CNCF)
  - *Motto: Building sustainable ecosystems for cloud native software*
  - *CNCF is part of the nonprofit Linux Foundation*
- Cloud Native = scalable apps running in modern cloud environments
  - *containers, service meshes, microservices*
  - *Apps must be usually re-built from scratch or refactored*
  - *Benefits:*
    - *loosely coupled systems that are resilient, manageable, and observable*
    - *automation allowing for predictable and frequent changes with minimal effort*
  - *Trail Map*
    - *provides an overview for enterprises starting their cloud native journey*
- Lift and Shift
  - *Cloud transition program in organizations*
  - *Move app from on-premise to the cloud*
  - *Benefits*
    - *Infrastructure cost cutting (OPEX vs. CAPEX)*
    - *Improved operations (scaling up/down if possible can be faster)*


# CNCF Trail Map

# Overview

- Cloud Native
- Kubernetes
  - *Basic Concepts*
  - *Core Concepts and Architecture*
  - *Workloads*
  - *Services*
  - *Beyond the Basics*

# Overview

- In your architecture...
  - Containers are atomic pieces of application architecture
  - Containers can be linked (e.g. web server, DB)
  - Containers access shared resources (e.g. disk volumes)
- Kubernetes
  - Automation of deployments, scaling, management of containerized applications across number of nodes
  - Based on Borg, a parent project from Google

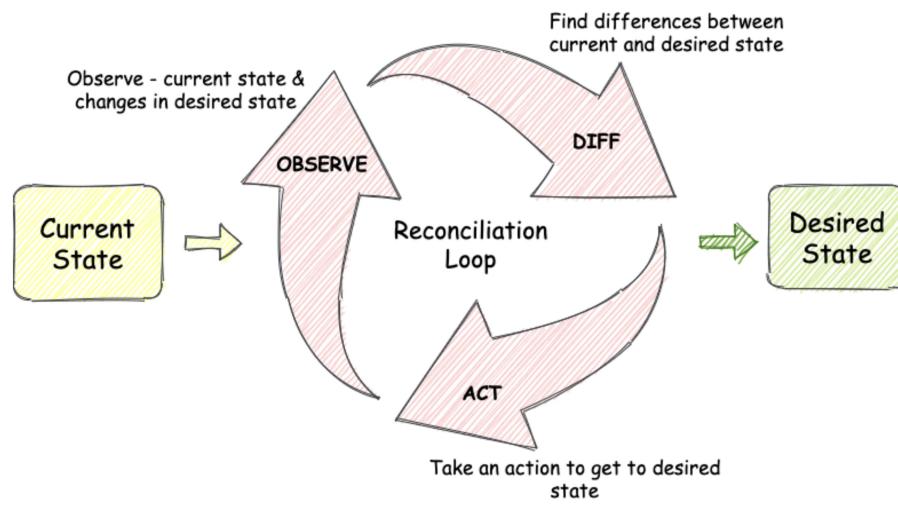


# Key Design Principles

- Kubernetes provides abstractions that separate application deployment from the underlying infrastructure details
- Application workloads and infrastructure decoupling
  - *Compute: Define what to run without specifying where it runs*
  - *Storage: Applications request storage independent of storage backend*
  - *Networking: Stable access to applications regardless of IPs or location*
- Benefits
  - *Portability across on-prem and cloud environments*
  - *Scalability and resilience through dynamic scheduling*
  - *Consistency and standardization of deployment model*
  - *Reduced vendor lock-in thanks to open standards*

# Desired State and Reconciliation

- Kubernetes operates on a **desired state** model
  - Users define the state they want through object specifications (YAML)
  - Example: “there should be 3 replicas of this application”
- Actual State vs. Desired State
  - Kubernetes constantly monitors the cluster
  - If the actual state drifts from the desired state, it takes action to fix it
- Reconciliation Loop
  - Controllers continuously compare desired vs. actual state
  - Automatically performs actions such as restarting, rescheduling, or scaling Pods



# Features

- Automatic binpacking
  - *Automatically places containers onto nodes based on their resource requirements and other constraints.*
- Horizontal scaling
  - *Scales your application up and down with a simple command, with a UI, or automatically based on CPU usage.*
- Automated rollouts and rollbacks
  - *Progressive rollout out of changes to application/configuration, monitoring application health and rollback when something goes wrong.*
- Storage orchestration
  - *Automatically mounts the storage system (local or in the cloud)*
- Self-healing
  - *Restarts containers that fail, replaces and reschedules containers when nodes die, kills containers that don't respond to user-defined health checks.*
- Service discovery and load balancing
  - *Gives containers their own IP addresses and a single DNS name for a set of containers, and can load-balance across them.*


# Overview

- Cloud Native
- Kubernetes
  - *Basic Concepts*
  - *Core Concepts and Architecture*
  - *Workloads*
  - *Services*
  - *Beyond the Basics*

# Core Building Blocks

- **Cluster**
  - *A set of worker nodes and a control plane*
  - *Runs and manages containerized applications*
- **Node**
  - *A worker machine in Kubernetes (VM or physical)*
  - *Runs Pods scheduled by the control plane*
- **Control Plane**
  - *Manages the overall state of the cluster*
  - *Schedules workloads and responds to cluster events*
- **Pod**
  - *The smallest deployable unit in Kubernetes*
  - *One or more tightly-coupled containers*
  - *Containers share networking and storage within a Pod*

# Architecture



# Control Plane Components (Part 1)

- Global decisions about the cluster
  - *Scheduling*
  - *Detecting and responding to cluster events, starting up new pods*
- kube-apiserver
  - *exposes the Kubernetes API*
  - *The API server is the front end for the Kubernetes control plane.*
- etcd
  - *highly-available key value store used to store all cluster data*
- kube-scheduler
  - *watches for newly created Pods with no assigned node*
  - *selects a node for Pods to run on.*
  - *Decision factors: resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications*


# Control Plane Components (Part 2)

- kube-controller-manager
  - *runs controller to ensure the desired state of cluster objects*
  - **Node controller**
    - *noticing and responding when nodes go down*
  - **Job controller**
    - *creates Pods to run one-off tasks to completion.*
  - **Endpoints controller**
    - *Populates the Endpoints object (that is, joins Services, Pods).*
- cloud-controller-manager
  - *Integration with cloud services (when the cluster is running in a cloud)*
  - **Node controller**
    - *checks if a node has been deleted in the cloud after it stops responding*
  - **Route controller**
    - *For setting up routes in the underlying cloud infrastructure*
  - **Service controller**

# Node

- Kubernetes runtime environment
  - *Run on every node*
  - *Maintaining running pods*
- kubelet
  - *An agent that runs on each node in the cluster*
  - *It makes sure that containers are running in a Pod.*
- kube-proxy
  - *maintains network rules on nodes*
  - *network rules allow network communication to Pods from inside or outside of the cluster*
  - *uses the operating system packet filtering layer or forwards the traffic itself.*
- Container runtime
  - *Responsible for running containers*
  - *Kubernetes supports several container runtimes (containerd, CRI-O)*
  - *Any implementation of the Kubernetes CRI (Container Runtime Interface)*

# Container Stack



# Overview

- Cloud Native
- Kubernetes
  - *Basic Concepts*
  - *Core Concepts and Architecture*
  - *Workloads*
  - *Services*
  - *Beyond the Basics*

# Namespaces

- Logical grouping of cluster resources
  - Allow you to organize and separate objects within a Kubernetes cluster
  - Useful when multiple teams, environments, or projects share the same cluster
- Rationale
  - Provide isolation and boundaries between workloads
  - Prevent name collisions
    - Objects can have the same name if in different namespaces
  - Enable resource limits and access control per namespace
- Usage
  - Common namespaces: **default**, **kube-system**, **kube-public**, **kube-node-lease**
  - Create separate namespaces for e.g. dev, test, prod
  - Commands run in a namespace unless another is specified

# Pod

- Pod
  - *A group of one or more tightly-coupled containers.*
  - *Containers share storage and network resources.*
  - *A Pod runs a single instance of a given application*
  - *Pod's containers are always co-located and co-scheduled*
  - *Pod's containers run in a shared context, i.e. in a set of Linux namespaces*
- Pods are created using workload resources
  - *You do not create them directly*
- Pods in a Kubernetes cluster are used in two main ways
  - *Run a single container, the most common Kubernetes use case*
  - *Run multiple containers that need to work together*

# Workloads

- An application running on Kubernetes
- Workloads run in a set of Pods
- Pre-defined workload resources to manage lifecycle of Pods
  - **Deployment** and **ReplicaSet**
    - managing a stateless application workload
    - any Pod in the Deployment is interchangeable and can be replaced if needed
  - **StatefulSet**
    - one or more related Pods that track state
    - For example, if a workload records data persistently, run a StatefulSet that matches each Pod with a persistent volume.
  - **DaemonSet**
    - Ensures that all (or some) Nodes run a copy of a Pod
    - Such as a cluster storage daemon, logs collection, node monitoring running on every node
  - **Job** and **CronJob**
    - Define tasks that run to completion and then stop.
    - Jobs represent one-off tasks, whereas CronJobs recur according to a schedule.

# Deployment Spec Example

- Deployment spec

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4      name: nginx-deployment
5  spec:
6      selector:
7          matchLabels:
8              app: nginx
9  replicas: 3 # tells deployment to run 3 pods matching the template
10 template:
11     metadata:
12         labels:
13             app: nginx
14     spec:
15         containers:
16             - name: nginx
17                 image: nginx:1.14.2
18             ports:
19                 - containerPort: 80
```

- *A desired state of an application running in the cluster*
- *Kubernetes reads the Deployment spec and starts three app instances*
- *If an instance fails, Kubernetes starts a replacement app instance*

# Overview

- Cloud Native
- Kubernetes
  - *Basic Concepts*
  - *Core Concepts and Architecture*
  - *Workloads*
  - *Services*
  - *Beyond the Basics*

# What is a Service?

- A Kubernetes **Service** is an abstraction that defines
  - *A logical set of Pods*
  - *A policy to access them.*
- Pods are ephemeral – their IPs change when recreated
- A Service provides a stable virtual endpoint for a set of Pods
- Services enable reliable communication between components:
  - *Internal pods communication*
  - *External access to cluster workloads*
- Each Service gets
  - *A DNS name and*
  - *virtual IP (ClusterIP) inside the cluster.*
- Kubernetes component **kube-proxy** manage routing to backend Pods.

# Service Types

- **ClusterIP**
  - Exposes the Service on an internal IP in the cluster only.
  - Used for internal communication between Pods.
- **NodePort**
  - Exposes the Service on each Node's IP at a static port (e.g. 30080).
  - Accessible externally via **NodeIP:NodePort**.
- **LoadBalancer**
  - Provisions an external load balancer (e.g. in cloud environments).
  - Routes external traffic to the Service.
- **ExternalName**
  - Maps the Service to an external DNS name.
  - No proxying — pure DNS CNAME redirection.

# How Services Work

- **Selector**
  - A Service usually defines a **selector** — a label query used to find matching Pods.
  - Example: **selector: app=nginx** matches all Pods with label **app=nginx**.
  - Kubernetes monitors Pods that match this selector and updates Service backends
- **Endpoints / EndpointSlice**
  - For every Service with a selector, Kubernetes creates an **Endpoints** (or **EndpointSlice**) object listing all healthy Pod IPs and ports.
  - This list changes dynamically as Pods are added, removed, or become unhealthy.
- **kube-proxy**
  - Runs on every Node and watches Service and Endpoint objects.
  - Programs **iptables** or **IPVS** rules to forward traffic from the Service's virtual IP (**ClusterIP**) to one of the backend Pod IPs.
  - Load balancing is done using round-robin or IPVS algorithms.
- **DNS Integration**
  - **CoreDNS** automatically creates a DNS record for each Service:
    - **<service>.⟨namespace⟩.svc.cluster.local**
  - Pods can reach the Service via DNS without knowing Pod IPs
    - **curl http://my-service.default.svc.cluster.local**

# ClusterIP Service Example

- Example configuration exposing an NGINX Deployment internally:

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: nginx-svc
5 spec:
6   selector:
7     app: nginx
8   ports:
9     - protocol: TCP
10    port: 80
11    targetPort: 8080
12  type: ClusterIP
```

- Pods with `app=nginx` receive traffic through **ClusterIP**.
- DNS name: `nginx-svc.default.svc.cluster.local`
- Used by other Pods to connect via `http://nginx-svc:80`.

# Packet Forwarding and Load Balancing

- **iptables mode**
  - *kube-proxy* creates NAT rules in the `nat` table to redirect Service traffic.
  - Example traffic comming to `NodeIP:NodePort` (e.g. `192.168.1.11:30080`)

```
1 # 1. Match NodePort traffic coming from outside
2 -A KUBE-NODEPORTS -p tcp --dport 30080 -m addrtype ! --src-type LOCAL \
3   -j KUBE-MARK-MASQ
4     # Mark all external traffic for SNAT (so replies go back via this node)
5
6 # 2. NodePort forwards traffic to the Service chain
7 -A KUBE-NODEPORTS -p tcp --dport 30080 \
8   -m comment --comment "default/my-service: NodePort" \
9   -j KUBE-SVC-XYZ123
10
11 # 3. Service chain chooses one backend Pod
12 -A KUBE-SVC-XYZ123 -m statistic --mode random --probability 0.5 \
13   -j KUBE-SEP-A1B2C3
14 -A KUBE-SVC-XYZ123 -j KUBE-SEP-D4E5F6
15
16 # 4. Pod DNAT rule to redirect to Pod IP:port
17 -A KUBE-SEP-A1B2C3 -p tcp -m tcp -j DNAT --to-destination 10.42.0.12:8080
18 -A KUBE-SEP-D4E5F6 -p tcp -m tcp -j DNAT --to-destination 10.42.1.7:8080
```
  - The node's routing table determines how to reach the Pod's IP:
    - `10.42.0.0/24` via `192.168.1.12 dev flannel.1`
    - packets to Pods in `10.42.0.0/24` (running on Node 2) are sent through the `VXLAN` interface `flannel.1` to Node 2's IP `192.168.1.12`

# Overview

- Cloud Native
- Kubernetes
  - *Basic Concepts*
  - *Core Concepts and Architecture*
  - *Workloads*
  - *Services*
  - *Beyond the Basics*

# Advanced Topics

- Custom APIs and Controllers
  - *CRDs, Operators, reconciliation loops*
  - *Admission webhooks (mutating/validating)*
- Security
  - *RBAC, Namespaces, Pod Security (seccomp, capabilities, rootless)*
  - *Image signing and supply chain (SBOM, cosign), Secret management (Vault/CSI)*
  - *Policy engines: OPA Gatekeeper, Kyverno*
- Networking
  - *CNI, eBPF (Cilium), NetworkPolicies, Ingress*
  - *Gateway API, Service Mesh (mTLS, traffic shaping)*
- Storage
  - *CSI drivers, snapshots, expansion, topology-aware PVs*
  - *Backup/DR (e.g., Velero), StatefulSet patterns*
- Scaling and Scheduling
  - *HPA/VPA/KEDA (event-driven), Cluster Autoscaler*
  - *Affinity/anti-affinity, taints/tolerations, topology spread*
- Ops and Delivery
  - *GitOps (Argo CD/Flux), progressive delivery (canary, blue/green)*