

Middleware Architectures 1

Lecture 5: Cloud Native and Kubernetes

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <https://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <https://vitvar.com/lectures>



Modified: Sun Oct 26 2025, 20:49:01
Humla v1.0

Overview

- Cloud Native
- Kubernetes

Overview

- The Cloud Native Computing Foundation (CNCF)
 - *Motto: Building sustainable ecosystems for cloud native software*
 - *CNCF is part of the nonprofit Linux Foundation*
- Cloud Native = scalable apps running in modern cloud environments
 - *containers, service meshes, microservices*
 - *Apps must be usually re-built from scratch or refactored*
 - *Benefits:*
 - *loosely coupled systems that are resilient, manageable, and observable*
 - *automation allowing for predictable and frequent changes with minimal effort*
 - *Trail Map*
 - *provides an overview for enterprises starting their cloud native journey*
- Lift and Shift
 - *Cloud transition program in organizations*
 - *Move app from on-premise to the cloud*
 - *Benefits*
 - *Infrastructure cost cutting (OPEX vs. CAPEX)*
 - *Improved operations (scaling up/down if possible can be faster)*

CNCF Trail Map




Overview

- Cloud Native
- Kubernetes
 - *Basic Concepts*
 - *Core Concepts and Architecture*
 - *Workloads*
 - *Beyond the Basics*

Overview

- In your architecture...
 - *Containers are atomic pieces of application architecture*
 - *Containers can be linked (e.g. web server, DB)*
 - *Containers access shared resources (e.g. disk volumes)*
- Kubernetes
 - *Automation of deployments, scaling, management of containerized applications across number of nodes*
 - *Based on Borg, a parent project from Google*




Key Design Principles

- Kubernetes provides abstractions that separate application deployment from the underlying infrastructure details
- Application workloads and infrastructure decoupling
 - **Compute:** Define what to run without specifying where it runs
 - **Storage:** Applications request storage independent of storage backend
 - **Networking:** Stable access to applications regardless of IPs or location
- Benefits
 - Portability across on-prem and cloud environments
 - Scalability and resilience through dynamic scheduling
 - Consistency and standardization of deployment model
 - Reduced vendor lock-in thanks to open standards

Desired State and Reconciliation

- Kubernetes operates on a **desired state** model
 - Users define the state they want through object specifications (YAML)
 - Example: “there should be 3 replicas of this application”
- Actual State vs. Desired State
 - Kubernetes constantly monitors the cluster
 - If the actual state drifts from the desired state, it takes action to fix it
- Reconciliation Loop
 - Controllers continuously compare desired vs. actual state
 - Automatically performs actions such as restarting, rescheduling, or scaling Pods



Features

- Automatic binpacking
 - Automatically places containers onto nodes based on their resource requirements and other constraints.
- Horizontal scaling
 - Scales your application up and down with a simple command, with a UI, or automatically based on CPU usage.
- Automated rollouts and rollbacks
 - Progressive rollout out of changes to application/configuration, monitoring application health and rollback when something goes wrong.
- Storage orchestration
 - Automatically mounts the storage system (local or in the cloud)
- Self-healing
 - Restarts containers that fail, replaces and reschedules containers when nodes die, kills containers that don't respond to user-defined health checks.
- Service discovery and load balancing
 - Gives containers their own IP addresses and a single DNS name for a set of containers, and can load-balance across them.


Overview

- Cloud Native
- Kubernetes
 - Basic Concepts
 - *Core Concepts and Architecture*
 - Workloads
 - Beyond the Basics

Core Building Blocks

- **Cluster**
 - A set of worker nodes and a control plane
 - Runs and manages containerized applications
- **Node**
 - A worker machine in Kubernetes (VM or physical)
 - Runs Pods scheduled by the control plane
- **Control Plane**
 - Manages the overall state of the cluster
 - Schedules workloads and responds to cluster events
- **Pod**
 - The smallest deployable unit in Kubernetes
 - One or more tightly-coupled containers
 - Containers share networking and storage within a Pod

Architecture



Control Plane Components (Part 1)

- Global decisions about the cluster
 - *Scheduling*
 - *Detecting and responding to cluster events, starting up new pods*
- kube-apiserver
 - *exposes the Kubernetes API*
 - *The API server is the front end for the Kubernetes control plane.*
- etcd
 - *highly-available key value store used to store all cluster data*
- kube-scheduler
 - *watches for newly created Pods with no assigned node*
 - *selects a node for Pods to run on.*
 - *Decision factors: resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications*


Control Plane Components (Part 2)

- kube-controller-manager
 - *runs controller to ensure the desired state of cluster objects*
 - **Node controller**
 - *noticing and responding when nodes go down*
 - **Job controller**
 - *creates Pods to run one-off tasks to completion.*
 - **Endpoints controller**
 - *Populates the Endpoints object (that is, joins Services, Pods).*
- cloud-controller-manager
 - *Integration with cloud services (when the cluster is running in a cloud)*
 - **Node controller**
 - *checks if a node has been deleted in the cloud after it stops responding*
 - **Route controller**
 - *For setting up routes in the underlying cloud infrastructure*
 - **Service controller**
 - *For creating, updating and deleting cloud provider load balancers*

Node

- Kubernetes runtime environment
 - *Run on every node*
 - *Maintaining running pods*
- kubelet
 - *An agent that runs on each node in the cluster*
 - *It makes sure that containers are running in a Pod.*
- kube-proxy
 - *maintains network rules on nodes*
 - *network rules allow network communication to Pods from inside or outside of the cluster*
 - *uses the operating system packet filtering layer or forwards the traffic itself.*
- Container runtime
 - *Responsible for running containers*
 - *Kubernetes supports several container runtimes (containerd, CRI-O)*
 - *Any implementation of the Kubernetes CRI (Container Runtime Interface)*

Container Stack



Overview

- Cloud Native
- Kubernetes
 - *Basic Concepts*
 - *Core Concepts and Architecture*
 - *Workloads*
 - *Beyond the Basics*

Namespaces

- Logical grouping of cluster resources
 - *Allow you to organize and separate objects within a Kubernetes cluster*
 - *Useful when multiple teams, environments, or projects share the same cluster*
- Rationale
 - *Provide isolation and boundaries between workloads*
 - *Prevent name collisions*
 - *Objects can have the same name if in different namespaces*
 - *Enable resource limits and access control per namespace*
- Usage
 - *Common namespaces: default, kube-system, kube-public, kube-node-lease*
 - *Create separate namespaces for e.g. dev, test, prod*
 - *Commands run in a namespace unless another is specified*

Pod

- Pod
 - A group of one or more tightly-coupled containers.
 - Containers share storage and network resources.
 - A Pod runs a single instance of a given application
 - Pod's containers are always co-located and co-scheduled
 - Pod's containers run in a shared context, i.e. in a set of Linux namespaces
- Pods are created using workload resources
 - You do not create them directly
- Pods in a Kubernetes cluster are used in two main ways
 - Run a single container; the most common Kubernetes use case
 - Run multiple containers that need to work together

Workloads

- An application running on Kubernetes
- Workloads run in a set of Pods
- Pre-defined workload resources to manage lifecycle of Pods
 - **Deployment** and **ReplicaSet**
 - managing a stateless application workload
 - any Pod in the Deployment is interchangeable and can be replaced if needed
 - **StatefulSet**
 - one or more related Pods that track state
 - For example, if a workload records data persistently, run a StatefulSet that matches each Pod with a persistent volume.
 - **DaemonSet**
 - Ensures that all (or some) Nodes run a copy of a Pod
 - Such as a cluster storage daemon, logs collection, node monitoring running on every node
 - **Job** and **CronJob**
 - Define tasks that run to completion and then stop.
 - Jobs represent one-off tasks, whereas CronJobs recur according to a schedule.

Deployment Spec Example

- Deployment spec

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5  spec:
6    selector:
7      matchLabels:
8        app: nginx
9    replicas: 3 # tells deployment to run 3 pods matching the template
10   template:
11     metadata:
12       labels:
13         app: nginx
14     spec:
15       containers:
16         - name: nginx
17           image: nginx:1.14.2
18           ports:
19             - containerPort: 80
```

- A desired state of an application running in the cluster
- Kubernetes reads the Deployment spec and starts three app instances
- If an instance fails, Kubernetes starts a replacement app instance

Service

- Networking

- Containers within a Pod use networking to communicate via loopback
- Cluster networking provides communication between different Pods.

- Service resource

- An abstract way to expose an application running on a set of Pods
- Example: a set of Pods with a label **app=nginx**, each listens on **tcp/9376**

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: my-service
5  spec:
6    selector:
7      app: nginx
8    ports:
9      - protocol: TCP
10        port: 80
11        targetPort: 9376
```

- This specification creates a new Service object named **my-service**
- The service targets **tcp/9376** on any Pod with the **app=nginx** label.
- Kubernetes assigns this Service a cluster IP address, which is used by the Service proxies.

Overview

- Cloud Native
- Kubernetes
 - *Basic Concepts*
 - *Core Concepts and Architecture*
 - *Workloads*
 - *Beyond the Basics*

Advanced Topics

- Custom APIs and Controllers
 - *CRDs, Operators, reconciliation loops*
 - *Admission webhooks (mutating/validating)*
- Security
 - *RBAC, Namespaces, Pod Security (seccomp, capabilities, rootless)*
 - *Image signing and supply chain (SBOM, cosign), Secret management (Vault/CSI)*
 - *Policy engines: OPA Gatekeeper, Kyverno*
- Networking
 - *CNI, eBPF (Cilium), NetworkPolicies, Ingress*
 - *Gateway API, Service Mesh (mTLS, traffic shaping)*
- Storage
 - *CSI drivers, snapshots, expansion, topology-aware PVs*
 - *Backup/DR (e.g., Velero), StatefulSet patterns*
- Scaling and Scheduling
 - *HPA/VPA/KEDA (event-driven), Cluster Autoscaler*
 - *Affinity/anti-affinity, taints/tolerations, topology spread*
- Ops and Delivery