

Middleware and Web Services

Lecture 6: High Availability and Performance

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <http://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <http://vitvar.com/courses/mdw>



Modified: Tue Nov 22 2016, 20:46:32
Humla v0.3

Good Performance

- What influences a good performance?
 - *Number of users and concurrent connections*
 - *Number of messages and messages' sizes*
 - *Number of services*
 - *Infrastructure – capacity, availability, configuration, ...*
- How can we achieve a good performance?
 - *Infrastructure*
 - *Scalability, failover, cluster architectures*
 - *Performance tuning*
 - *Application Server, JVM memory, OS-level tuning, Work managers configuration*
 - *Service configuration*
 - *Parallel processing, process optimization*

Overview

- **Infrastructure**
 - *Load Balancers*
 - *Cluster Architecture*
- **Performance Tuning**

Definitions

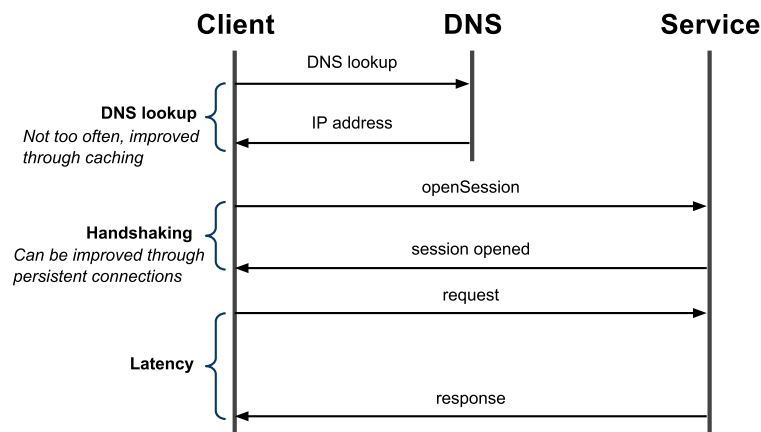
- **Scalability**
 - *server scalability*
 - *ability of a system to scale – when input load changes*
 - *users should not feel a difference when more users access the same application at the same time*
 - **horizontal scaling**
 - *adding new instances of applications/servers*
 - **vertical scaling**
 - *adding new resources (CPU, memory) to a server instance*
 - *network traffic*
 - *bandwidth capacity influences performance too*
 - *service should limit the network traffic through caching*
- **Availability**
 - *probability that a service is operational at a particular time*
 - *e.g., 99.9987% availability – downtime ~44 seconds/year*

Definitions (Cont.)

- **High Availability**
 - When a server instance fails, operation of the application can continue
 - Failures should affect application availability and performance as little as possible
- **Application Failover**
 - When an application component performing a job becomes unavailable, a copy of the failed object finishes the job.
 - Issues
 - A copy of the failed object must be available
 - A location and operational status of available objects must be available
 - A processing state must be replicated
- **Load Balancing**
 - Distribution of incoming requests across server instances

Performance Metrics

- **Latency**
 - A client-side metric

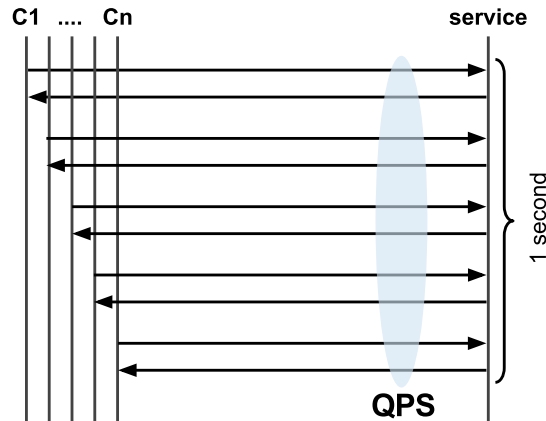


- CPU intensive service or a bad configuration of a service
 - consider asynchronous processing when CPU intensive
- Writing to a data store

Performance Metrics

- Queries/Requests per Second (QPS)

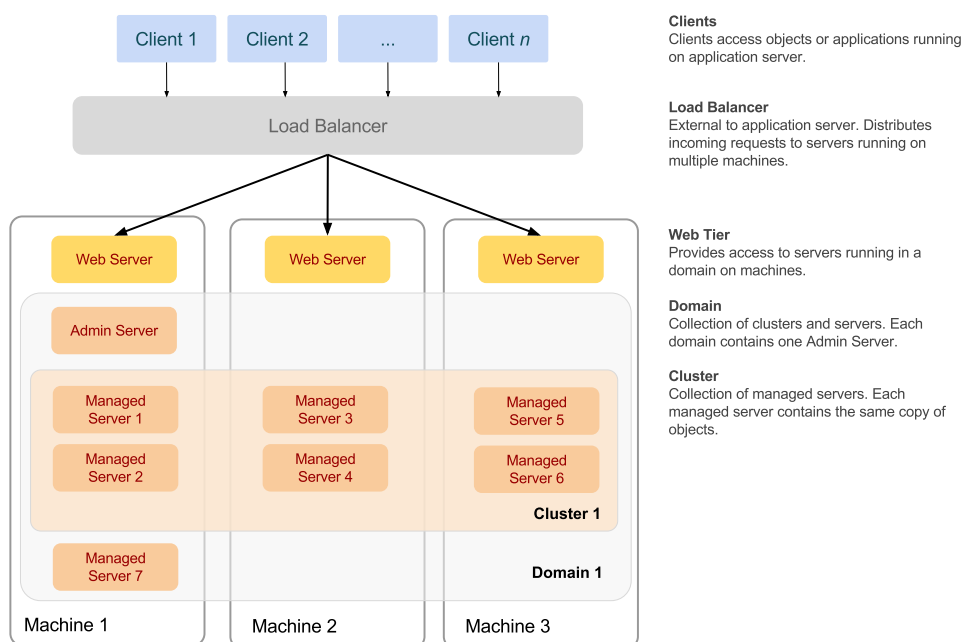
– *A server-side metric*



– *Caching may improve performance*

→ *even if data changes often, with high QPS caching improves a lot*

Infrastructure Example – Weblogic



Best Configuration Practices

- Domain configuration
 - *A server is an admin server or a managed server*
 - *Each server is running on a separated JVM*
 - *A physical machine may run one or more servers*
 - *There should be at least two managed servers running on one machine*
 - *This provides a better performance*
(as opposed to one server running on one machine)
 - *A domain can have clustered or unclustered servers*
- Load balancers (LB)
 - *Load Balancers are not part of the domain*
 - *They are external to Weblogic server*
 - *There is usually one HW LB and several SW LBs*
 - *Software LB*
 - *Realized by the Web Tier (Apache HTTP server)*
 - *Redirects requests too all managed servers in a domain (across multiple machines)*

Overview

- Infrastructure
 - *Load Balancers*
 - *Cluster Architecture*
- Performance Tuning

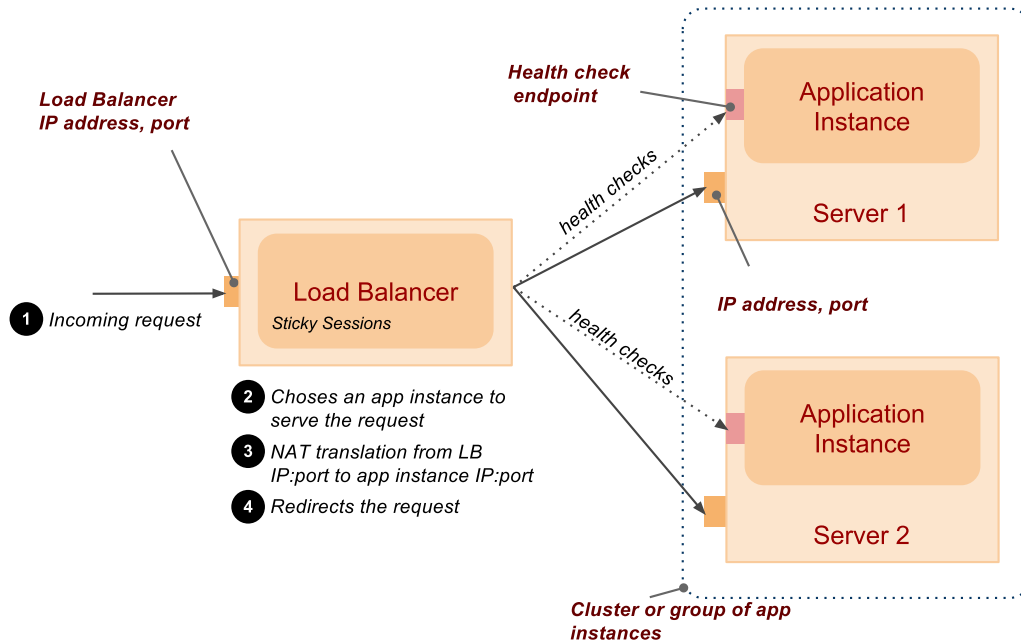
Load Balancing

- Distributes a load to multiple app/object instances
 - App instances run on different machines
 - Load sharing: equal or with preferences
 - Health checks
- Types
 - DNS-based load balancer
 - DNS Round Robin
 - NAT-based load balancer (Layer-4)
 - **Reverse-proxy load balancer** (Layer-7)
 - application layer
 - Sticky sessions
 - JSession, JSession-aware load balancer
 - Client-side load balancer
 - LB run by a client
 - a client uses a replica-aware stub of the object from the server

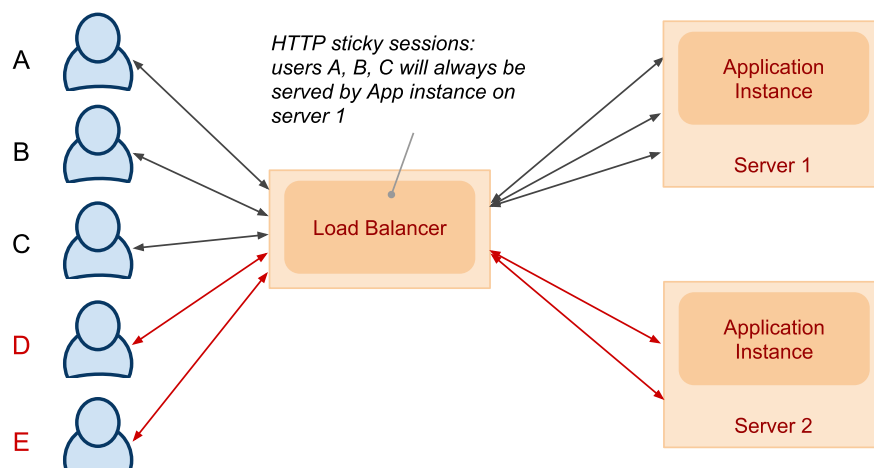
DNS-based Load Balancer

- DNS Round Robin
 - A DNS record has multiple assigned IP addresses
 - DNS system delivers different IP addresses from the list
 - Example DNS A Record:
`company.com A 147.32.100.71 147.32.100.72 147.32.100.73`
- Advantages
 - Very simple, easy to implement
- Disadvantages
 - IP address in cache, could take hours to re-assign
 - No information about servers' loads and health

Reverse Proxy Load Balancer



HTTP Sticky Sessions Example



- How to identify a server that hosts the session state
 - Passive cookie persistence – LB uses a cookie from the server
 - Active cookie persistence – LB adds its own cookie

Types of Load Balancers

- Software

- *Apache mod_proxy_balancer*
 - *HTTP Session persistence – sticky sessions*
- *WebLogic proxy plug-in*

```
1 <Location /soa-infra>
2   SetHandler weblogic-handler
3   WebLogicCluster czfmwapp03-vf:8001,czfmwapp04-vf:8001,czfmwapp05-vf:800
4 </Location>
5
```

/soa-infra is a first part of an URL path that rules in this **Location** will be applied (this is a standard Apache configuration mechanism)

czfmwapp{N} is a hostname that corresponds to a virtual IP to which the managed server JVM processes is bounded (using the tcp port **8001**).

WebLogicCluster specifies the list of servers for load balancing

- Hardware

- *Cisco, Avaya, Barracude*

Round-Robin Algorithm with Health Check

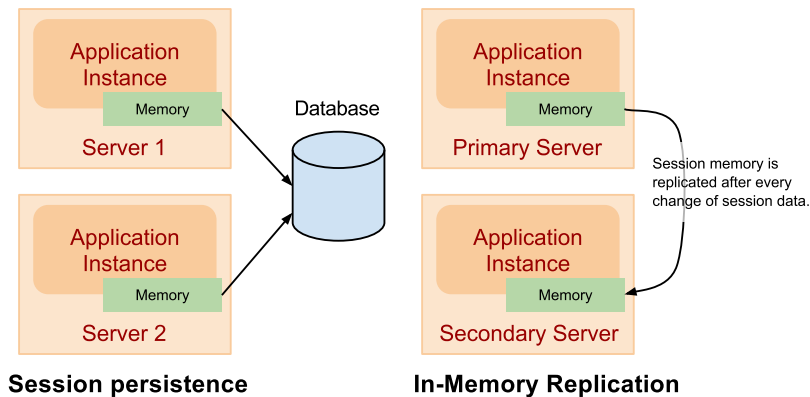
- Uses

- *request – client request with or without a cookie information*
- *server list – a list of servers that can process the request*
 - e.g. **WebLogicCluster** value (see previous slide)
- *unhealthy treshhold – a number of negative consecutive health checks before moving the server to the "unhealthy" state.*

- Steps

- *if a cookie exist in the request that identifies a server*
- *always use that server*
- **health check**
- *LB polls the servers' healthcheck endpoints*
- *if a number of health checks exceeds the unhealthy threshold*
 - *LB removes the server from the server list*
- *if a server was unhealthy and a there was a successful healthcheck*
 - *LB adds the server to the server list*

Session State Persistence and Replication



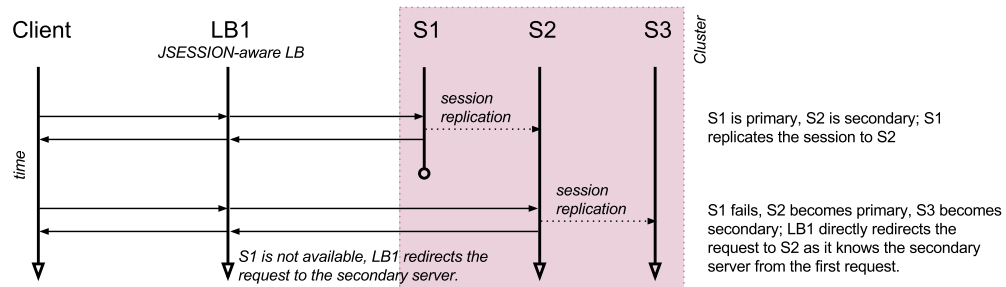
- Session persistence
 - Session information is maintained in the database
 - Does not require sticky sessions
 - Implements `HttpSession` interface that writes data to the DB
- In-memory replication
 - A **primary server** holds a session state, the **secondary server** holds its replica.
 - Information about primary and secondary servers are part of `JSession`

In-Memory Replication

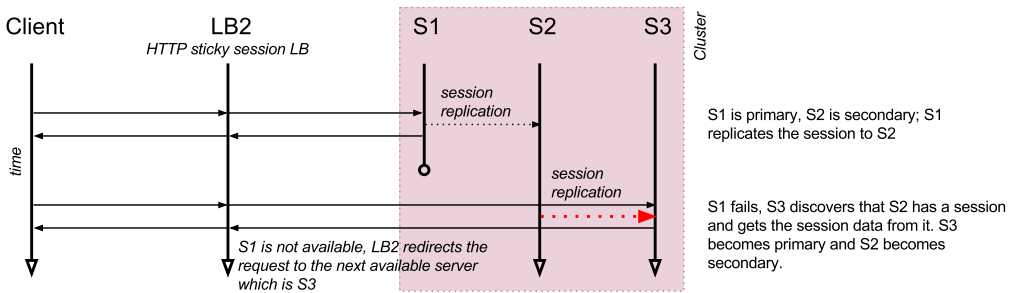
- Session format
 - It's a cookie
 - `JSESSIONID=SESSION_ID!PRIM_SERVER_ID!SEC_SERVER_ID!CREATION_TIME`
 - `SESSION_ID` – session id, generated by the server to identify memory associated with the session on the server
 - `PRIM_SERVER_ID` – ID of the managed server holding the session data
 - `SEC_SERVER_ID` – ID of the managed server holding the session replica
 - `CREATION_TIME` – time the session data was created/updated
- How LB uses this information
 - LB has information whether the server is running or not (via healthchecks)
 - if the primary server is running, it redirects the request there
 - if the primary server is not running, it redirects the request to the secondary server directly
 - if primary and secondary servers are not running, it redirect the request to any other server it has in the list – this may cause side effects!

In-Memory Replication Scenarios

Scenario A: JSession-aware load balancer



Scenario B: HTTP sticky session load balancer



Overview

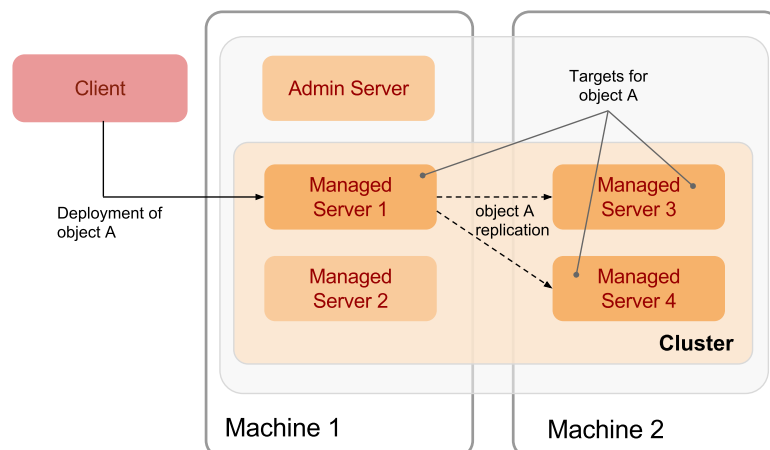
- Infrastructure
 - Load Balancers
 - Cluster Architecture
- Performance Tuning

Overview

- Cluster
 - A group of servers that act together to serve client requests
 - Cluster appears to clients as a single application server
 - Servers can run on the same machines or on different machines
 - Cluster's capacity can be increased by adding servers to the cluster
 - Servers in a cluster may have the same copy of objects and they are aware of each other objects
 - objects: applications, JMS destinations, RMI objects
- Communication in the cluster
 - peer-to-peer communication using IP sockets
 - IP multicast which servers use to broadcast availability of objects and heartbeats
- Configurations
 - Objects deployed to all servers in a cluster
 - Cluster-wide JNDI tree allows to look-up clustered objects
 - Servers in a cluster may get replicated through migration

Deployment to Cluster

- Deployment of an object
 - Client deploys to one managed server in the cluster
 - Object gets replicated to its targets
 - Targets can be configured for the object, usually all servers but can be selected servers
 - See [Lecture 4](#) for the definition of the object

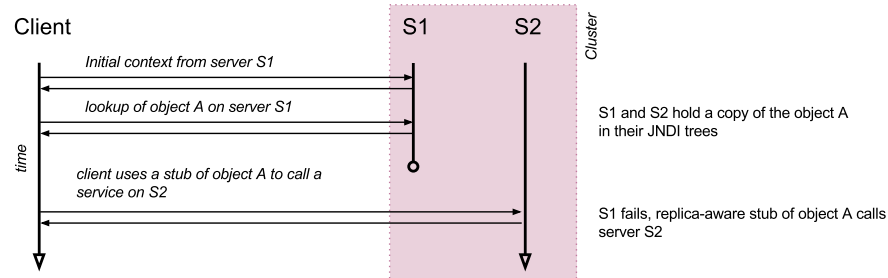


Object Failover

- Failover

- *Failover = ability to locate an object on another server that holds a copy of the object without an impact on the performance and configuration*

Replica-aware stub of object A, failover in cluster

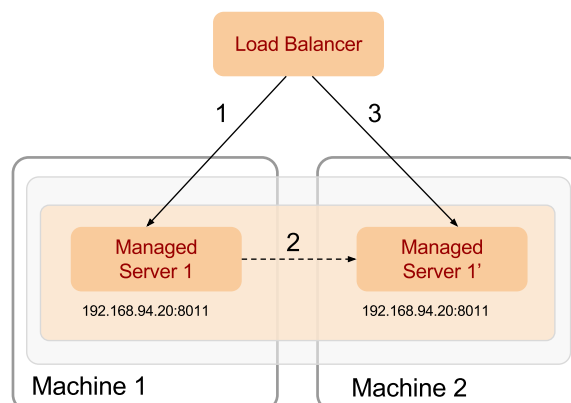


- *A client gets a stub of the object by calling **lookup** on the context*
- *A client uses the stub of the object to access the object on the server*
- *When a server fails, replicate-aware stub calls the next server that holds the object copy*

Server Failover

- Failover

- *Failover = ability to relocate the server to another machine without an impact on the performance*



- *Managed server listens on **virtual_IP:port***
- *A load balancer forwards a request to **virtual_IP:port***
- *When the server moves, **virtual_IP:port** remains the same*

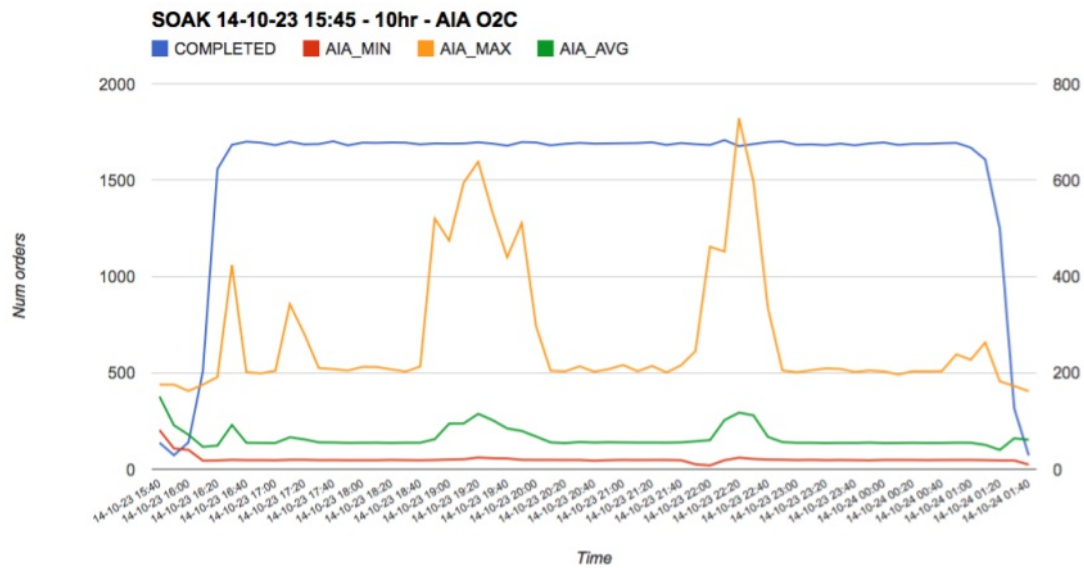
Overview

- Infrastructure
- Performance Tuning

Performance Limiting Factors



Example Performance Testing

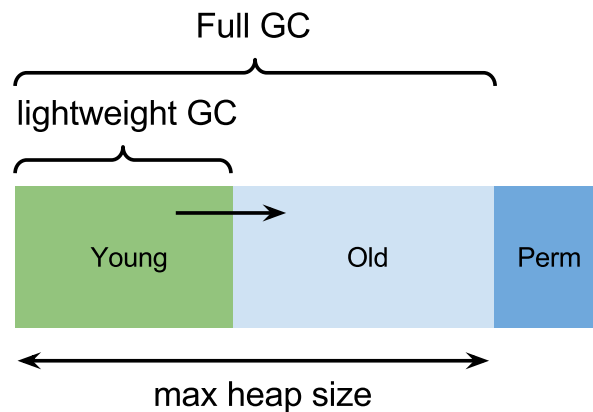


- Completed – number of completed orders
- MIN, MAX, AVG – a minimum/maximum/average processing time in 10 minutes
- At 18:30–20:20 was a performance issue with OMS environment

Tuning – A Layered Approach

- Application server can be tuned at multiple layers
 - Service configuration optimization
 - Transport-level tuning
 - Application Server Tuning
 - JVM Memory Tuning
 - OS Tuning
- Lower levels are cheaper to tune

Memory Allocations



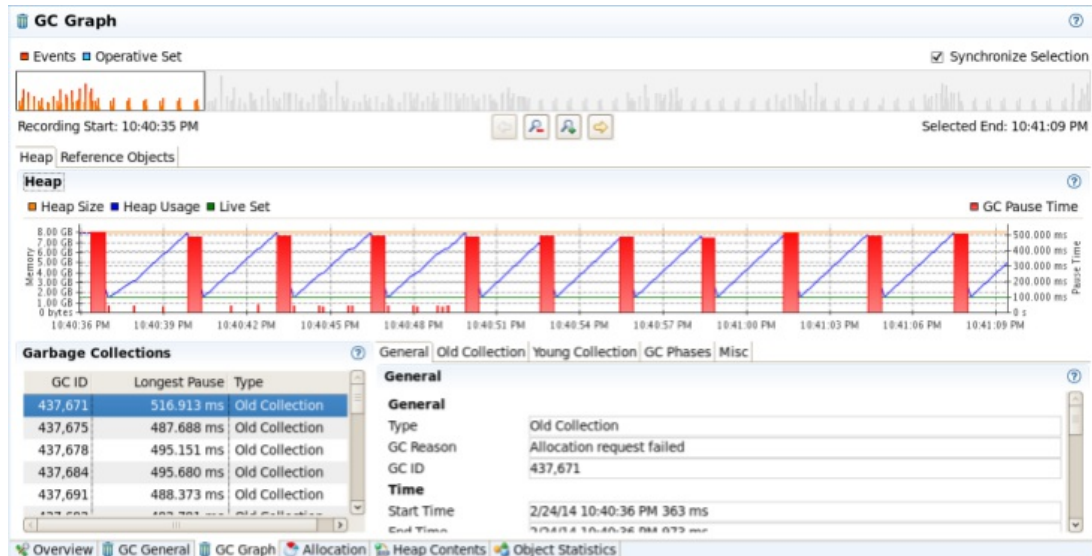
- Generations
 - Young – objects get allocated in this space initially
 - Old – objects get promoted to old from young
 - Perm – space for permanent allocations, e.g. objects describing classes and methods

Garbage Collection

- Steps to move objects around
 1. Objects are created in young
 2. When young is full, the live objects are copied to old, dead are discarded
 - **lightweight GC**
 3. When young is full and no space in old → the full GC frees the old space
 - **Full GC** – nothing is running in JVM, the application stops
 - **Too frequent full GC has an impact on performance**
- A memory leak or inadequate heap allocation
 - Old is out of space → full GC will run often (or continuously)
 - High CPU utilization, ESB will not be able to process/respond to requests

Heap Size and GC Runs

- Heap Size and GC runs
 - *Wrong heap size allocation – too small or memory leaks*
 - *GC full runs too often, this has a negative impact on performance*

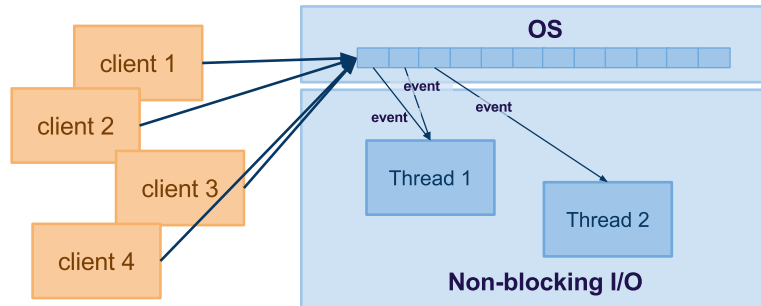


JVM Memory Tuning

- JVM Memory Parameters
 - **-Xms** – *initial java heap size*
 - **-Xmx** – *maximum java heap size*
 - **-XX:NewSize** – *the initial size of the heap for young generation*
 - **-XX:MaxNewSize** – *the maximum size of the heap for young generation*
- General recommendations
 - **-Xms** and **-Xmx** should be set to the same value
(do not allow the heap to grow → limit the overhead)
 - **-XX:NewSize** and **-XX:MaxNewSize** should be set to the one half of maximum heap
 - Example, 1GB heap size
 - **-Xms1024m -Xmx1024m -XX:NewSize=500m -XX:MaxNewSize=500m**

Asynchronous I/O: Recall

- Connections maintained by the OS, not the Web app
 - *The Web app registers events, OS triggers events when occur*



- Characteristics
 - *Event examples: new connection, read, write, closed*
 - *The app may create working threads, but controls the number!*
 - *much less number of working threads as opposed to blocking I/O*

Work Manager Configuration

- Work Manager
 - *Controls the number of thread allocated to processing of requests*
 - *In WLS is called a dispatch policy*
 - *Can be assigned to OSB proxy services*
 - *Parameters*
 - **maximum threads (max)** – *maximum number of working threads*
 - **capacity (cap)** – *maximum number of connections*
 - *maximum connections waiting to be processed: **cap - max***
 - *refused connections: when number of connections is **> cap***
- Inbound throttling
 - *A dispatch policy applied to a single proxy service*
 - *Rejected connections will not be processed*