

# Middleware and Web Services

## Lecture 6: High Availability and Performance

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <http://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <http://vitvar.com/courses/mdw>



Modified: Sun Nov 19 2017, 23:02:59  
Humla v0.3

## Good Performance

- What influences good performance?
  - *Number of users and concurrent connections*
  - *Number of messages and messages' sizes*
  - *Number of services*
  - *Infrastructure – capacity, availability, configuration, ...*
- How can we achieve good performance?
  - *Infrastructure*
    - *Scalability, failover, cluster architectures*
  - *Performance tuning*
    - *Application Server, JVM memory, OS-level tuning, Work managers configuration*
  - *Service configuration*
    - *Parallel processing, process optimization*

## Overview

- **Infrastructure**
  - *Load Balancers*
  - *Cluster Architecture*
- **Performance Tuning**

## Definitions

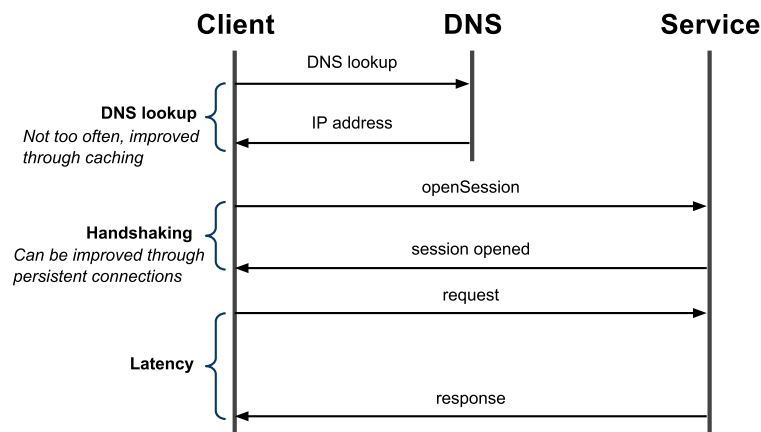
- **Scalability**
  - *server scalability*
    - *ability of a system to scale – when input load changes*
    - *users should not feel a difference when more users access the same application at the same time*
    - **horizontal scaling**
      - *adding new instances of applications/servers*
    - **vertical scaling**
      - *adding new resources (CPU, memory) to a server instance*
  - *network traffic*
    - *bandwidth capacity influences performance too*
    - *service should limit the network traffic through caching*
- **Availability**
  - *probability that a service is operational at a particular time*
    - *e.g., 99.9987% availability – downtime ~44 seconds/year*

## Definitions (Cont.)

- **High Availability**
  - When a server instance fails, operation of the application can continue
  - Failures should affect application availability and performance as little as possible
- **Application Failover**
  - When an application component performing a job becomes unavailable, a copy of the failed object finishes the job.
  - Issues
    - A copy of the failed object must be available
    - A location and operational status of available objects must be available
    - A processing state must be replicated
- **Load Balancing**
  - Distribution of incoming requests across server instances

## Performance Metrics

- **Latency**
  - A client-side metric

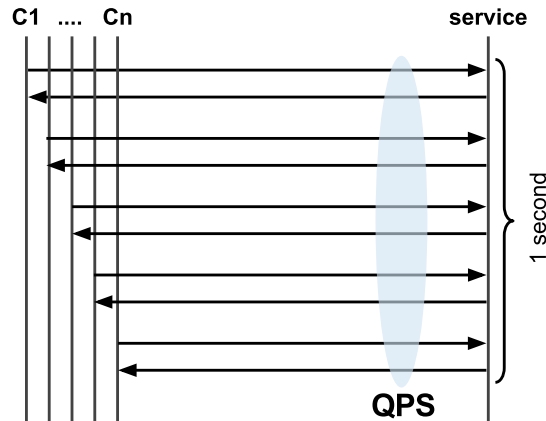


- CPU intensive service or a bad configuration of a service
  - consider asynchronous processing when CPU intensive
- Writing to a data store

## Performance Metrics

- Queries/Requests per Second (QPS)

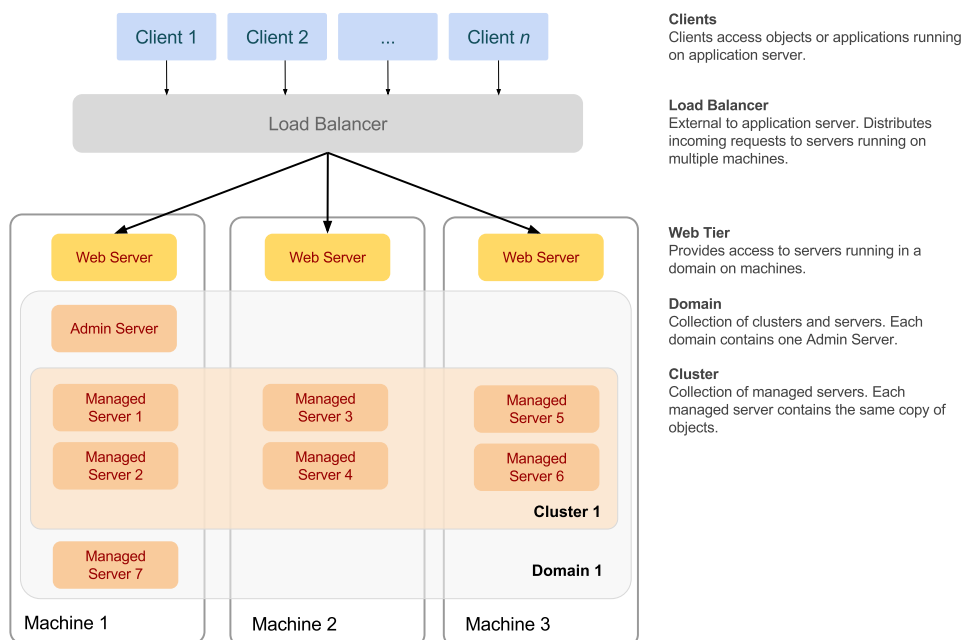
– *A server-side metric*



– *Caching may improve performance*

→ *even if data changes often, with high QPS caching improves a lot*

## Infrastructure Example – Weblogic



## Best Configuration Practices

- Domain configuration
  - *A server is an admin server or a managed server*
  - *Each server is running on a separated JVM*
  - *A physical machine may run one or more servers*
  - *There should be at least two managed servers running on one machine*
    - *This provides a better performance*  
*(as opposed to one server running on one machine)*
  - *A domain can have clustered or unclustered servers*
- Load balancers (LB)
  - *Load Balancers are not part of the domain*
    - *They are external to Weblogic server*
    - *There is usually one HW LB and several SW LBs*
    - *Software LB*
      - *Realized by the Web Tier (Apache HTTP server)*
      - *Redirects requests too all managed servers in a domain (across multiple machines)*

## Overview

- Infrastructure
  - *Load Balancers*
  - *Cluster Architecture*
- Performance Tuning

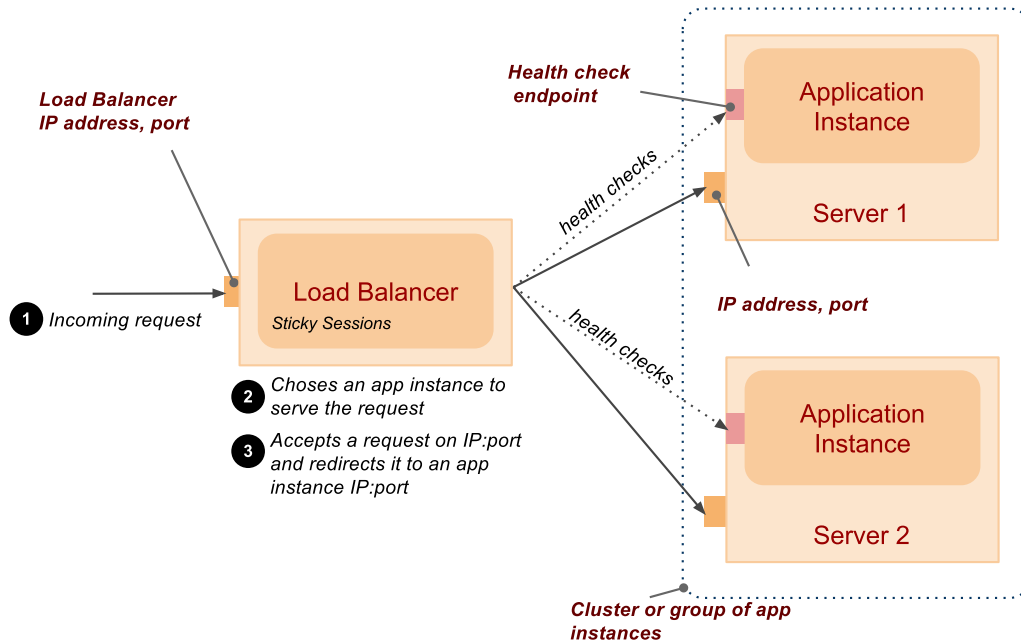
## Load Balancing

- Distributes a load to multiple app/object instances
  - App instances run on different machines
  - Load sharing: equal or with preferences
  - Health checks
- Types
  - DNS-based load balancer
    - DNS Round Robin
  - NAT-based load balancer (Layer-4)
  - **Reverse-proxy load balancer** (Layer-7)
    - application layer
    - Sticky sessions
      - JSession, JSession-aware load balancer
  - Client-side load balancer
    - LB run by a client
    - a client uses a replica-aware stub of the object from the server

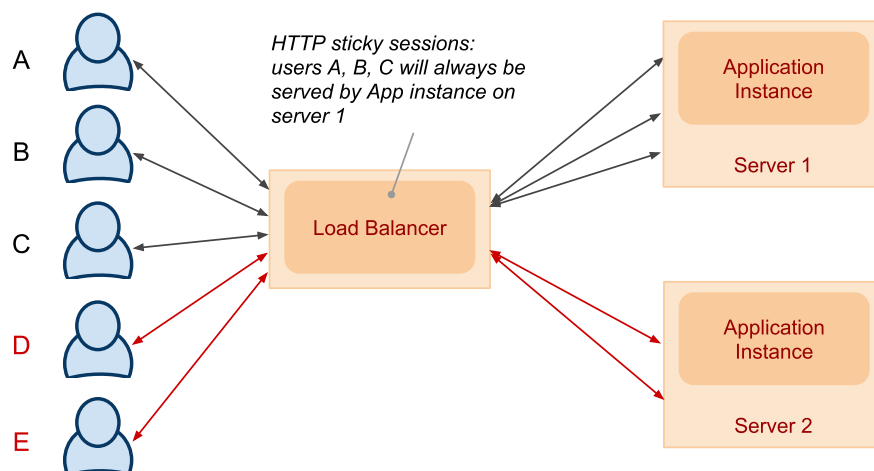
## DNS-based Load Balancer

- DNS Round Robin
  - A DNS record has multiple assigned IP addresses
  - DNS system delivers different IP addresses from the list
  - Example DNS A Record:  
`company.com A 147.32.100.71 147.32.100.72 147.32.100.73`
- Advantages
  - Very simple, easy to implement
- Disadvantages
  - IP address in cache, could take hours to re-assign
  - No information about servers' loads and health

# Reverse Proxy Load Balancer



# HTTP Sticky Sessions Example



- How to identify a server that hosts the session state
  - *Passive cookie persistence* – LB uses a cookie from the server
  - *Active cookie persistence* – LB adds its own cookie

## Types of Load Balancers

- Software

- *Apache mod\_proxy\_balancer*
  - *HTTP Session persistence – sticky sessions*
- *WebLogic proxy plug-in*

```
1 <Location /soa-infra>
2   SetHandler weblogic-handler
3   WebLogicCluster czfmwapp03-vf:8001,czfmwapp04-vf:8001,czfmwapp05-vf:800
4 </Location>
5
```

*/soa-infra* is a first part of an URL path that rules in this **Location** will be applied (this is a standard Apache configuration mechanism)

**czfmwapp{N}** is a hostname that corresponds to a virtual IP to which the managed server JVM processes is bounded (using the tcp port **8001**).

**WebLogicCluster** specifies the list of servers for load balancing

- Hardware

- *Cisco, Avaya, Barracude*

## Round-Robin Algorithm

- Uses

**request** – client request with or without a cookie information

**server\_list** – a list of servers that can process the request

**rbinx** – round robin index

**sticky\_sessions** – associative array of pairs **<session\_id,server>**

**unhealthy\_treshold** – a number of negative consecutive health checks before moving the server to the "unhealthy" state.

- Round Robin Algorithm


- if **session\_id** exist in the **request** and in **sticky\_sessions**
  - send the **request** to the server **sticky\_sessions[session\_id]**
- otherwise
  - send the **request** to the **rbinx** server in the **server\_list**
  - extract **session\_id** from the response from the server
  - if the **session\_id** exist, add a pair **<session\_id;server\_list[rbinx]>** to **sticky\_sessions**
  - increase **rbinx** by one or reset it to **0** if it exceeds the length of **server\_list**



## Health Check

- Health Check
  - For each server in the `server_list`
    - call the server's healthcheck endpoint
    - if a number of failed health checks for the server exceeds the `unhealthy_threshold`
      - remove the server from the `server_list`
    - if the server was unhealthy and there was a successful healthcheck
      - add the server back to the `server_list`

## Session State Persistence and Replication



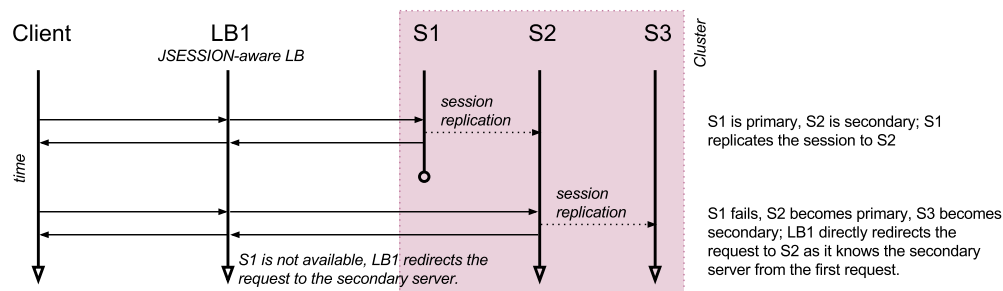
- Session persistence
  - Session information is maintained in the database
  - Does not require sticky sessions
  - Implements `HttpSession` interface that writes data to the DB
- In-memory replication
  - A **primary server** holds a session state, the **secondary server** holds its replica.
  - Information about primary and secondary servers are part of `JSession`

# In-Memory Replication


- Session format
  - It's a cookie
  - `JSESSIONID=SESSION_ID!PRIM_SERVER_ID!SEC_SERVER_ID!CREATION_TIME`
    - `SESSION_ID` – session id, generated by the server to identify memory associated with the session on the server
    - `PRIM_SERVER_ID` – ID of the managed server holding the session data
    - `SEC_SERVER_ID` – ID of the managed server holding the session replica
    - `CREATION_TIME` – time the session data was created/updated
- How LB uses this information
  - LB has information whether the server is running or not (via healthchecks)
  - if the primary server is running, it redirects the request there
  - if the primary server is not running, it redirects the request to the secondary server directly
  - if primary and secondary servers are not running, it redirect the request to any other server it has in the list – this may cause side effects!

## In-Memory Replication Scenarios

### Scenario A: JSession-aware load balancer



### Scenario B: HTTP sticky session load balancer



## Overview


- Infrastructure
  - *Load Balancers*
  - *Cluster Architecture*
- Performance Tuning

## Overview

- Cluster capabilities
  - *A group of servers (aka cluster members) act together to serve clients' requests*
  - *Cluster is transparent to clients*
  - *Servers can run on the same machines or on different machines*
  - *Cluster's capacity can be increased by adding servers to the cluster*
  - *Servers in a cluster may have the same copy of objects and they are aware of each other objects*
    - *objects: applications, JMS destinations, RMI objects*
    - *See [Cluster-wide JNDI tree](#) in Lecture 4*
- Cluster Messaging Protocols
  - *When servers need to send messages to other members of the cluster*
  - **Unicast** - *one-to-one communication using TCP/IP sockets*
  - **Multicast** - *one-to-many communication*
  - *Cluster services that rely on the cluster messaging protocol*
    - *Cluster Membership*
    - *JNDI Replication*
  - *There are services using persistent RMI connections between cluster members*
    - *such as during in-memory replication of HTTP session information*

## Unicast

- Overview
  - Uses standard TCP/IP sockets to send messages across cluster members
  - Uses a **group leader strategy** to limit the number of sockets required to send messages within the cluster
- Groups in the cluster
  - Cluster is split into **M** groups with **N** servers, each group has a leader
  - Servers communicate with the group leader, group leaders communicate with servers in the group and other group leaders
  - When a group leader dies, the group elects another group leader
  - There is up to **MxN** network messages for every cluster message



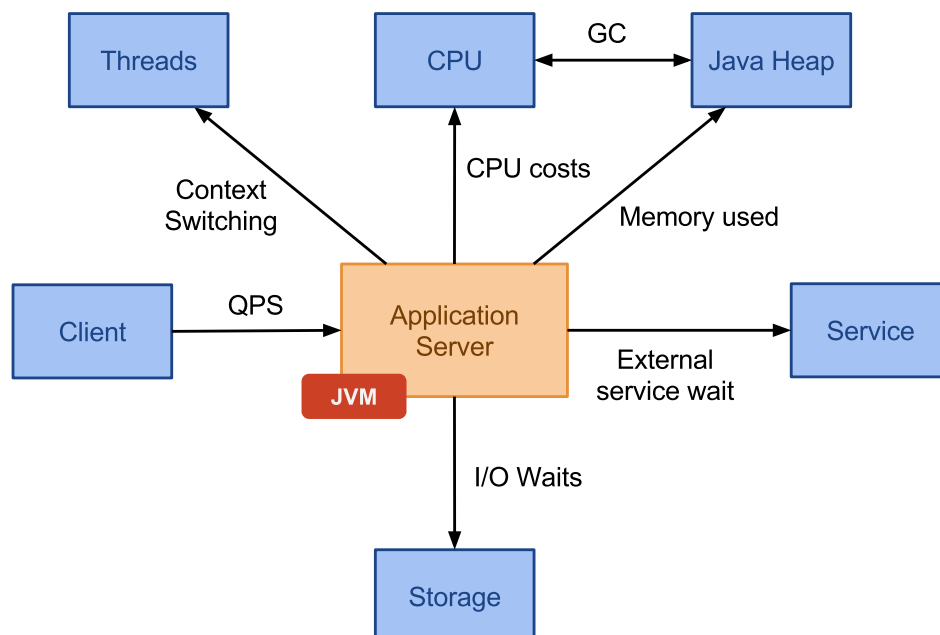
## Cluster Services

- Cluster Membership
  - Cluster members maintain their view on what servers are in the cluster
  - Each server sends heartbeat messages to the cluster
    - others know it is alive
  - Each server receives heartbeat messages
    - When a server misses a number of messages from a cluster, it removes the server from the list until it receives the next heartbeat message.
- JNDI Replication
  - Provides each server with a cluster-wide view of the JNDI tree.
  - Servers send JNDI update messages to the cluster when an object is bound or removed to their local tree.
  - When a server leaves the cluster
    - other members remove its JNDI bindings from their tree.
  - When a server joins the cluster
    - The server asks other server for a JNDI state dump.

## Overview

- Infrastructure
- Performance Tuning

## Performance Limiting Factors



## Monitoring

- Important to understand performance
  - *DevOps monitoring trends*
- What you need
  - *Collect → Filter → Store → View → **Tune***
  - *Metrics, dashboards, alerting, log management, reporting, tracing capabilities*
  - *It is necessary to organize metrics well in order to understand what is going on*
  - *Start from a high-level process, detail to technical components*
- Source
  - *Application server*
    - *usually management beans with JMX interfaces*
    - *log files (access logs, server logs, etc.)*
  - *OS*
    - *many utilities available out of the box*
    - *open sockets, memory, context switches, I/O performance, CPU usage*
  - *Database*
    - *applications may write metrics to the DB*
    - *SQL scripts to collect metrics*


## Monitoring Tools

- Commercial Monitoring Solutions
  - *Application server vendor usually offers a monitoring solution*
  - *AppDynamics, Oracle Enterprise Manager, Splunk*
  - *Google stackdriver, Amazon AWS CloudWatch*
- Open source examples
  - *Elasticsearch + LogStash + Kibana*
  - *InfluxDB + Telegraph + DataGraph*

## Tuning – A Layered Approach

- Application server can be tuned at multiple layers
  - *Service configuration optimization*
  - *Transport-level tuning*
  - *Application Server Tuning*
  - *JVM Memory Tuning*
  - *OS Tuning*
- Lower levels are cheaper to tune

## Memory Allocations



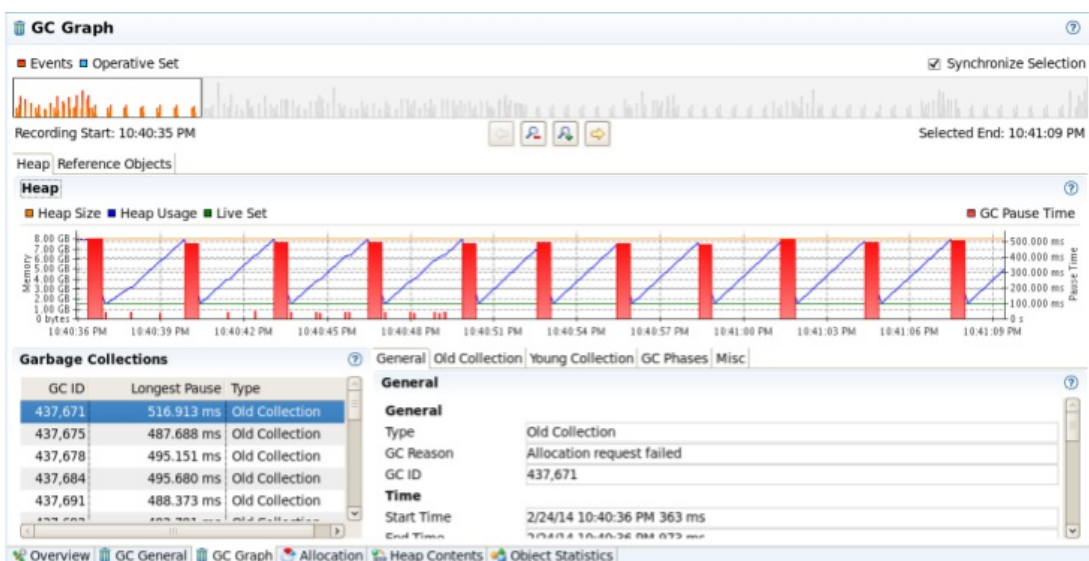
- Generations
  - *Young* – objects get allocated in this space initially
  - *Old* – objects get promoted to old from young
  - *Perm* – space for permanent allocations, e.g. objects describing classes and methods

# Garbage Collection

- Steps to move objects around
  1. Objects are created in young
  2. When young is full, the live objects are copied to old, dead are discarded – **lightweight GC**
  3. When young is full and no space in old → the full GC frees the old space – **Full GC** – nothing is running in JVM, the application stops – **Too frequent full GC has an impact on performance**
- A memory leak or inadequate heap allocation
  - Old is out of space → full GC will run often (or continuously)
  - High CPU utilization, The server will not be able to process/respond to requests

## Heap Size and GC Runs

- Heap Size and GC runs
  - Wrong heap size allocation – too small or memory leaks
  - GC full runs too often, this has a negative impact on performance






## JVM Memory Tuning

- JVM Memory Parameters
  - Xms – initial java heap size
  - Xmx – maximum java heap size
  - XX:NewSize – the initial size of the heap for young generation
  - XX:MaxNewSize – the maximum size of the heap for young generation
- General recommendations
  - Xms and -Xmx should be set to the same value  
(do not allow the heap to grow → limit the overhead)
  - XX:NewSize and -XX:MaxNewSize should be set to the one half of maximum heap
  - Example, 1GB heap size
    - Xms1024m -Xmx1024m -XX:NewSize=500m -XX:MaxNewSize=500m

## Asynchronous I/O: Recall

- Connections maintained by the OS, not the Web app
  - The Web app registers events, OS triggers events when occur



- Characteristics
  - Event examples: new connection, read, write, closed
  - The app may create working threads, but controls the number!  
→ much less number of working threads as opposed to blocking I/O

## Work Manager Configuration

- Work Manager
  - Controls the number of thread allocated to processing of requests
  - In WLS is called a dispatch policy
    - Can be assigned to OSB proxy services
  - Parameters
    - **maximum threads (max)** – maximum number of working threads
    - **capacity (cap)** – maximum number of connections
  - maximum connections waiting to be processed: **cap - max**
  - refused connections: when number of connections is > **cap**
- Inbound throttling
  - A dispatch policy applied to a single proxy service
  - Rejected connections will not be processed