Middleware Architectures 1
Lecture 7: HATEOAS, Caching and Concurrency

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • https://vitvar.com

Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • https://vitvar.com/lectures

---

Overview

- HATEOAS
- Caching, Revalidation, Concurrency Control
- Richardson Maturiy Model

HATEOAS

- HATEOAS = Hypertext as the Engine for Application State
  - *The REST core principle*
  - ***Hypertext***
    - → *Hypertext is a representation of a resource state with **links***
    - → *A link is an URI of a resource*
    - → *Applying an access (PUT, POST, DELETE) to a resource via its link = state transition*
- Statelessness
  - *A service does not use a session memory to remember a state*
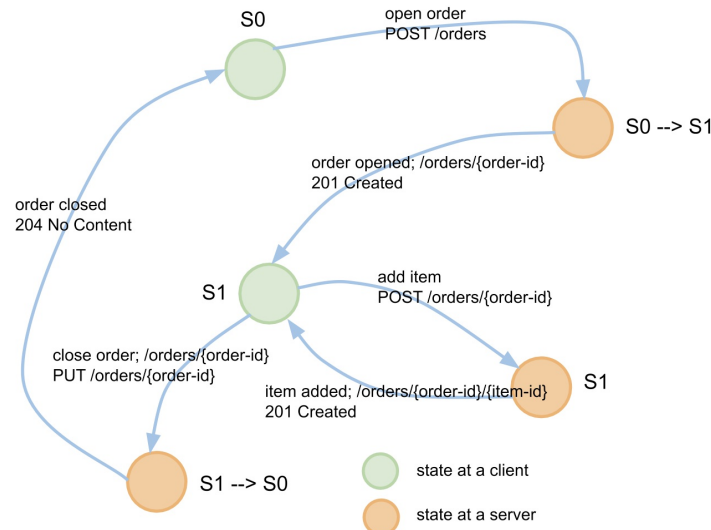 - *HATEOAS enables stateless implementation of services*

Stateful server

- Sessions to store the application state
  - *The app uses a server memory to remember the state*
  - *When the server restarts, the app state is lost*

Stateless server

- HTTP and hypermedia to transfer the app state
  - *Does not use a server memory to remember the app state*
 - *State transferred between a client and a service via HTTP metadata and resources' representations*

S0

open order
POST /orders

S0 --> S1

order opened; /orders/{order-id}
201 Created

order closed
204 No Content

S1

add item
POST /orders/{order-id}

close order; /orders/{order-id}
PUT /orders/{order-id}

item added; /orders/{order-id}/{item-id}
201 Created

S1

S1 --> S0

state at a client

state at a server

Persistent Storage and Session Memory

- Persistent Storage
  - *Contains the app data*
  - *Data is serialized into resource representation formats*
  - *All sessions may access the data via resource IDs*

- Session Memory
  - *Server memory that contains a state of the app*
  - *A session may only access its session memory*
  - *Access through cookies*
  - *Note*
    - → *A session memory may be implemented via a persistent storage (such as in Google AppEngine)*

Link

- Service operation
  - *Applying an access to a link (GET, PUT, POST, DELETE)*
  - *Link: HTTP method + resource URI + optional link semantics*

- Example: `getOrder`, `addOrder`, and `updateItem`
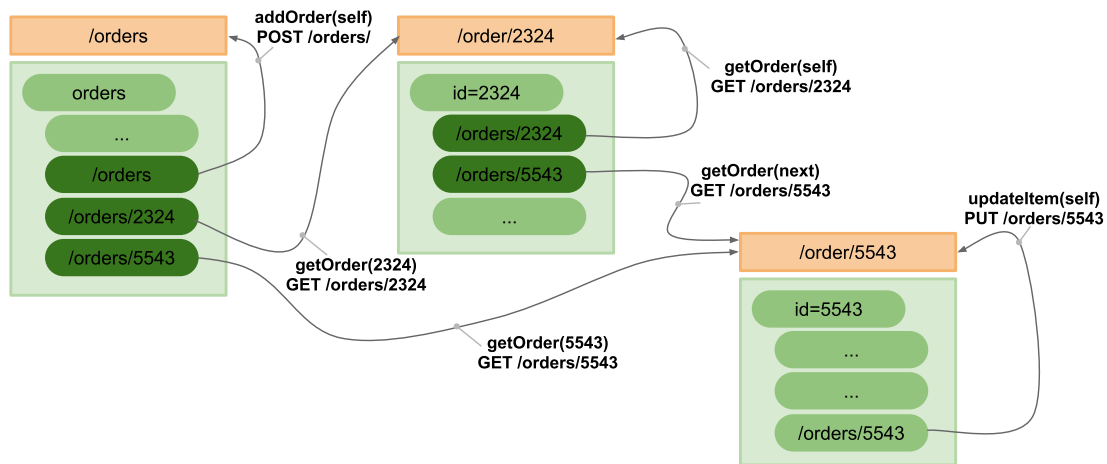
Atom Links

- Atom Syndication Format
  - *XML-based document format; Atom feeds*
  - *Atom links becoming popular for RESTful applications*

```
1  <order a:xmlns="http://www.w3.org/2005/Atom" xmlns="...">
2      <a:link
3          rel="next"
4          href="http://company.com/orders/5543"
5          type="application/xml"/>
6      <customer>Tomas</customer>
7      <items>...</items>
8  </order>
```

  - *Link structure*

    `rel` – *name of the link*

    ~ *semantics of an operation behind the link*

    `href` – *URI to the resource described by the link*

    `type` – *media type of the resource the link points to*

Link Semantics

- Standard `rel` values
  - *Navigation:* `next`, `previous`, `self`
  - *Does not reflect a HTTP method you can use*
- Extension `rel` values
  - *You can use `rel` to indicate a semantics of an operation*
  - *Example: add item, delete order, update order, etc.*
  - *A client associates this semantics with an operation it may apply at a particular state*
  - *The semantics should be defined by using an URI*

```
1  <order a:xmlns="http://www.w3.org/2005/Atom" xmlns="...">
2      <id>2324</id>
3      <a:link rel="http://company.com/op/addItem"
4          href="http://company.com/orders/2324"/>
5      <a:link rel="http://company.com/op/deleteOrder"
6          href="http://company.com/orders/2324"/>
7  </order>
```

Link Headers

- An alternative to Atom links in resource representations
  - *links defined in HTTP Link header, Web Linking IETF spec*
  - *They have the same semantics as Atom Links*
  - *Example:*

```
> HEAD /orders HTTP/1.1

< Content-Type: application/xml
< Link: <http://company.com/orders/?page=2&size=10>; rel="next"
< Link: <http://company.com/orders/?page=10&size=10>; rel="last"
```
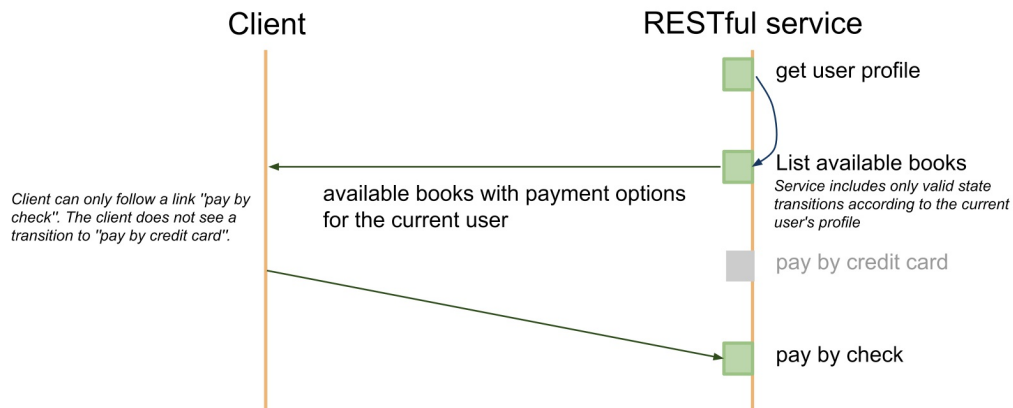
- Advantages
  - *no need to get the entire document*
  - *no need to parse the document to retrieve links*
  - *use HTTP HEAD only*

Preconditions and HATEOAS

- Preconditions in HATEOAS
  - *Service in a current state generates only valid transitions that it includes in the representation of the resource.*
  - *Transition logic is realized at the server-side*

Client   RESTful service

get user profile

List available books
Service includes only valid state transitions according to the current user's profile

Client can only follow a link "pay by check". The client does not see a transition to "pay by credit card".

available books with payment options for the current user

pay by credit card

pay by check

Advantages

- Location transparency
  - *only "entry-level" links published to the World*
  - *other links within documents can change without changing client's logic*
  - *Hypertext represents the current user's view, i.e. rights or other context*

- Loose coupling
  - *no need for a logic to construct the links*
  - *Clients know to which states they can move via links*

- Statelessness and Cloud
  - *Better implementation of scalability*

Overview

- HATEOAS
- Caching, Revalidation, Concurrency Control
- Richardson Maturiy Model

Scalability

- Need for scalability
  - *Huge amount of requests on the Web every day*
  - *Huge amount of data downloaded*
- Some examples
  - *Google, Facebook: 5 billion API calls/day*
  - *Twitter: 3 billions of API calls/day (75% of all the traffic) → 50 million tweets a day*
  - *eBay: 8 billion API calls/month*
  - *Bing: 3 billion API calls/month*
  - *Amazon WS: over 100 billion objects stored in S3*
- Scalability in REST
  - *Caching and revalidation*
  - *Concurrency control*

Caching

- ① request resource (GET)
- ② Cache-Control header
- ③ Revalidation

Client — Proxy — Server

Local and proxy cache to cache responses from the server

server cache to cache storage access and remote resource access.

- Your service should cache:
  - *anytime there is a static resource*
  - *even there is a dynamic resource*
    - → *with chances it updates often*
    - → *you can force clients to always revalidate*
- three steps:
  - *client GETs the resource representation*
  - *server controls how it should cache through* `Cache-Control` *header*
  - *client revalidates the content via conditional GET*

Cache Headers

- `Cache-Control` response header
  - *controls over local and proxy caches*
  - `private` – *no proxy should cache, only clients can*
  - `public` – *any intermediary can cache (proxies and clients)*
  - `no-cache` – *the response should not be cached. If it is cached, the content should always be revalidated.*
  - `no-store` – *must not store persistently (this turns off caching)*
  - `no-transform` – *no transformation of cached data; e.g. compressions*
  - `max-age`, `s-maxage` *a time in seconds how long the cache is valid;* `s-maxage` *for proxies*
- `Last-Modified` and `ETag` response headers
  - *Content last modified date and a content entity tag*
- `If-Modified-Since` and `If-None-Match` request headers
  - *Content revalidation (conditional GET)*

Example Date Revalidation

- Cache control example:

```
> GET /orders HTTP/1.1
> ...

< HTTP/1.1 200 OK
< Content-Type: application/xml
< Cache-Control: private, max-age=200
< Last-Modified: Sun, 7 Nov 2011, 09:40 CET
<
< ...data...
```

  – *only client can cache, the cache is valid for 200 seconds.*

- Revalidation (conditional GET) example:
  – *A client revalidates the cache after* **200** *seconds.*

```
> GET /orders HTTP/1.1
> If-Modified-Since: Sun, 7 Nov 2011, 09:40 CET

< HTTP/1.1 304 Not Modified
< Cache-Control: private, max-age=200
< Last-Modified: Sun, 7 Nov 2011, 09:40 CET
```

Entity Tags

- Signature of the response body
  – *A hash such as MD5*
  – *A sequence number that changes with any modification of the content*

- Types of tag
  – *Strong ETag: reflects the content bit by bit*
  – *Weak ETag: reflects the content "semantically"*
    → *The app defines the meaning of its weak tags*

- Example content revalidation with `ETag`

```
< HTTP/1.1 200 OK
< Cache-Control: private, max-age=200
< Last-Modified: Sun, 7 Nov 2011, 09:40 CET
< ETag: "4354a5f6423b43a54d"

> GET /orders HTTP/1.1
> If-None-Match: "4354a5f6423b43a54d"

< HTTP/1.1 304 Not Modified
< Cache-Control: private, max-age=200
< Last-Modified: Sun, 7 Nov 2011, 09:40 CET
< ETag: "4354a5f6423b43a54d"
```

Design Suggestions

- Composed resources use weak ETags
  - *For example* `/orders`
    - → *a composed resource that contains a summary information*
    - → *changes to an order's items will not change semantics of* `/orders`
  - *It is usually not possible to perform updates on these resources*
- Non-composed resources use strong ETags
  - *For example* `/orders/{order-id}`
  - *They can be updated*
- Further notes
  - *Server should send both* `Last-Modified` *and* `ETag` *headers*
  - *If client sends both* `If-Modified-Since` *and* `If-None-Match`, `ETag` *validation takes preference*

Weak ETag Example

- App specific, `/orders` resource example

```
1  {
2      "orders" :
3          [
4              { "id"       : 2245,
5                "customer" : "Tomas",
6                "descr"    : "Stuff to build a house.",
7                "items"    : [...] },
8              { "id"       : 5546,
9                "customer" : "Peter",
10               "descr"    : "Things to build a pipeline.",
11               "items"    : [...] }
12         ]
13 }
```

- Weak ETag compute function example
  - *Any modification to an order's items is not significant for* `/orders`:

```
1  var crypto = require("crypto");
2
3  function computeWeakETag(orders) {
4      var content = "";
5      for (var i = 0; i < orders.length; i++)
6          content += orders[i].id + orders[i].customer + orders[i].descr;
7      return crypto.createHash('md5').update(content).digest("hex");
8  }
```

Weak ETag Revalidation

- Updating `/orders` resource
  - `POST /orders/{order-id}` *inserts a new item to an order*
  - *Any changes to orders' items will not change the Weak ETag*


```
Client 1                                    Resource    Resource              Client 2
                                            /orders     /orders/{order-id}

GET /orders HTTP/1.1

HTTP/1.1 200 OK
Cache-Control: public, max-age=200
Last-Modified: Sun, 7 Nov 2011, 09:40 CET
ETag: W/"43ef54a4e32f80a43ef54b32"

                                                        POST /orders/2245 HTTP/1.1

>200                                                    HTTP/1.1 201 Created
seconds
GET /orders/ HTTP/1.1
If-None-Match:W/"43ef54a4e32f80a43ef54b32"

HTTP/1.1 304 Not Modified
Cache-Control: public, max-age=200
Last-Modified: Sun, 7 Nov 2011, 09:40 CET
ETag: W/"43ef54a4e32f80a43ef54b32"
```

Concurrency

- Two clients may update the same resource
  - *1) a client GETs a resource* `GET /orders/5545`
  - *2) the client modifies the resource*
  - *3) the client updates the resource via* `PUT /orders/5545 HTTP/1.1`

  What happens if another client updates the resource between 1) and 3) ?

- Concurrency control
  - *Conditional* `PUT`
    - → *Update the resource only if it has not changed since a specified date or a specified ETag matches the resource content*
  - `If-Unmodified-Since` *and* `If-Match` *headers*
  - *Response to conditional* `PUT`:
    - → `200 OK` *if the* `PUT` *was successful*
    - → `412 Precondition Failed` *if the resource was updated in the meantime.*

Concurrency Control Protocol

Client 1 **Server** **Client 2**

(1) GET /orders/4456

(2) GET /orders/4456

(3) HTTP/1.1 200 OK
Cache-Control: public, max-age=200
Last-Modified: Sun, 7 Nov 2010, 09:40 CET
...

(4) HTTP/1.1 200 OK
Cache-Control: public, max-age=200
Last-Modified: Sun, 7 Nov 2010, 09:40 CET
...

(5) PUT /orders/4456
If-Unmodified-Since: Sun, 7 Nov 2010, 09:40 CET
updated data

(6) PUT /orders/4456
If-Unmodified-Since: Sun, 7 Nov 2010, 09:40 CET
updated data

(7) HTTP/1.1 200 OK

(8) HTTP/1.1 412 Precondition Failed

- Conditional PUT and ETags
  – *Conditional PUT must always use strong entity tags or date validation*

Overview

- HATEOAS
- Caching, Revalidation, Concurrency Control
- Richardson Maturiy Model

Steps towards REST

See Richardson Maturiy Model details.

Levels

- LEVEL 0 – POX (Plain Old XML)
  - *HTTP as a tunneling mechanism*
  - *URL defines a service endpoint*
  - *No Web principles*

- LEVEL 1 – Resources
  - *Take advantages of resources and URIs*

- LEVEL 2 – HTTP Verbs
  - *Use HTTP methods and respect their semantics*

- LEVEL 3 – Hypermedia Controls
  - *HATEOAS*