

# Middleware Architectures 1

## Lecture 6: High Availability and Performance

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <https://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <https://vitvar.com/lectures>



Modified: Sun Dec 05 2021, 22:09:30  
Humla v1.0

## Good Performance

- What influences good performance?
  - *Number of users and concurrent connections*
  - *Number of messages and messages' sizes*
  - *Number of services*
  - *Infrastructure – capacity, availability, configuration, ...*
- How can we achieve good performance?
  - *Infrastructure*
    - *Scalability, failover, cluster architectures*
  - *Performance tuning*
    - *Application Server, JVM memory, OS-level tuning, Work managers configuration*
  - *Service configuration*
    - *Parallel processing, process optimization*

## Overview

- Definitions
- Load Balancers
- Performance Tuning

## Definitions

- Scalability
  - *server scalability*
    - *ability of a system to scale – when input load changes*
    - *users should not feel a difference when more users access the same application at the same time*
    - **horizontal scaling**
      - *adding new instances of applications/servers*
    - **vertical scaling**
      - *adding new resources (CPU, memory) to a server instance*
  - *network traffic*
    - *bandwidth capacity influences performance too*
    - *service should limit the network traffic through caching*
- Availability
  - *probability that a service is operational at a particular time*
    - *e.g., 99.9987% availability – downtime ~44 seconds/year*
- SLA – Service Level Agreement
  - *Guarantee of service availability*
  - *When availability is below a guaranteed value, a customer can get a discount*

## Definitions (Cont.)

- **High Availability**
  - When a server instance fails, operation of the application can continue
  - Failures should affect application availability and performance as little as possible
- **Application Failover**
  - When an application component performing a job becomes unavailable, a copy of the failed object finishes the job.
  - Issues
    - A copy of the failed object must be available
    - A location and operational status of available objects must be available
    - A processing state must be replicated
- **Load Balancing**
  - Distribution of incoming requests across server instances

## Performance Metrics

- **Response Time**
  - A client-side metric



- CPU intensive service or a bad configuration of a service
  - consider asynchronous processing when CPU intensive
- Writing to a data store

## Performance Metrics

- Queries/Requests per Second (QPS)

- A server-side metric



- Caching may improve performance

- even if data changes often, with high QPS caching improves a lot

## Overview

- Definitions
- **Load Balancers**
- Performance Tuning

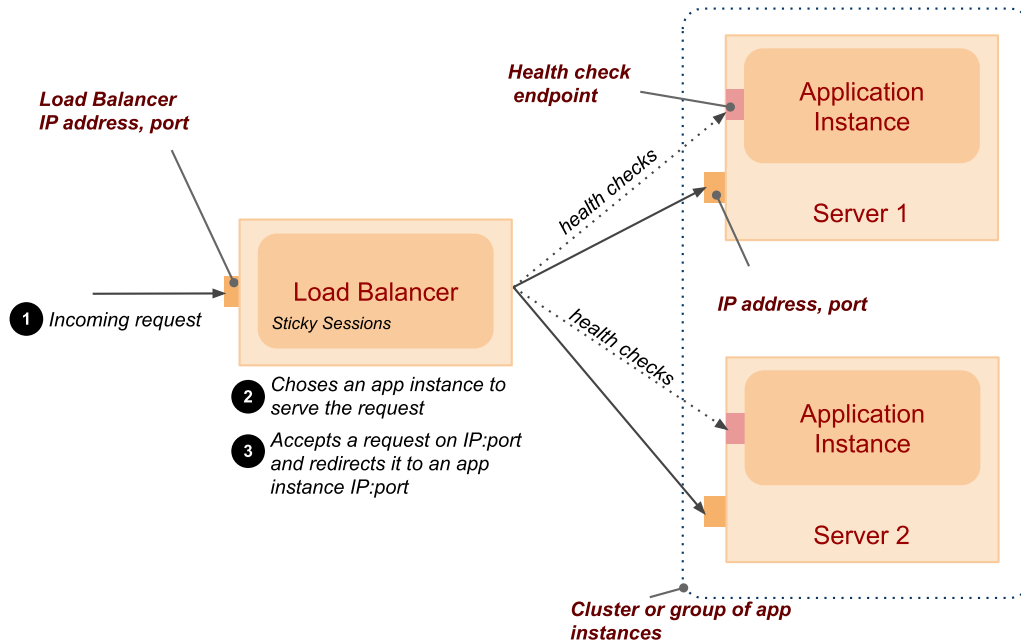
## Load Balancing

- Distributes a load to multiple app/object instances
  - App instances run on different machines
  - Load sharing: equal or with preferences
  - Health checks
- Types
  - DNS-based load balancer
    - DNS Round Robin
  - NAT-based load balancer (Layer-4)
  - **Reverse-proxy load balancer** (Layer-7)
    - application layer
    - Sticky sessions
      - JSession, JSession-aware load balancer
  - Client-side load balancer
    - LB run by a client
    - a client uses a replica-aware stub of the object from the server

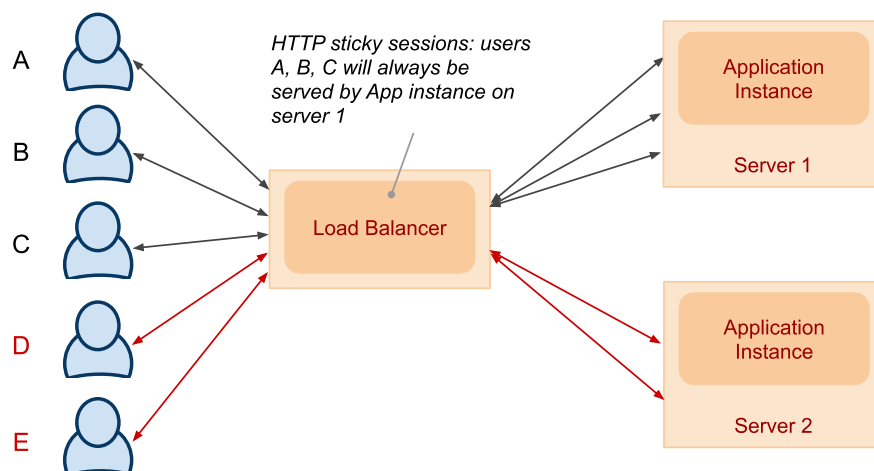
## DNS-based Load Balancer

- DNS Round Robin
  - A DNS record has multiple assigned IP addresses
  - DNS system delivers different IP addresses from the list
  - Example DNS A Record:  
`company.com A 147.32.100.71 147.32.100.72 147.32.100.73`
- Advantages
  - Very simple, easy to implement
- Disadvantages
  - IP address in cache, could take hours to re-assign
  - No information about servers' loads and health

# Reverse Proxy Load Balancer



# HTTP Sticky Sessions Example



- How to identify a server that hosts the session state
  - *Passive cookie persistence* – LB uses a cookie from the server
  - *Active cookie persistence* – LB adds its own cookie

## Types of Load Balancers

- Software

- Apache `mod_proxy_balancer`, `NGINX`
  - HTTP Session persistence – sticky sessions
  - Various configuration options
- WebLogic proxy plug-in

```
1 <Location /soa-infra>
2     SetHandler weblogic-handler
3     WebLogicCluster czfmwapp03-vf:8001,czfmwapp04-vf:8001,czfmwapp05-vf:800
4 </Location>
5
```

`/soa-infra` is a first part of an URL path that rules in this `Location` will be applied (this is a standard Apache configuration mechanism)  
`WebLogicCluster` specifies the list of backend servers for load balancing

- Hardware

- Cisco, Avaya, Barracude

## Round-Robin Algorithm

- Uses

`request` – client request with or without a cookie information  
`server_list` – a list of backend servers that can process the request  
`rbinx` – round robin index  
`sticky_sessions` – associative array of pairs `<session_id,server>`  
`unhealthy_treshold` – a number of negative consecutive health checks before moving the server to the "unhealthy" state.

- Round Robin Algorithm

- if `session_id` exist in the `request` and in `sticky_sessions`
  - send the `request` to the server `sticky_sessions[session_id]`
- otherwise
  - send the `request` to the `rbinx` server in the `server_list`
  - extract `session_id` from the response from the server
  - if the `session_id` exist, add a pair `<session_id;server_list[rbinx]>` to `sticky_sessions`
  - increase `rbinx` by one or reset it to 0 if it exceeds the length of `server_list`

## Health Check

- Health Check
  - For each server in the `server_list`
    - call the server's healthcheck endpoint
    - if a number of failed health checks for the server exceeds the `unhealthy_threshold`
      - remove the server from the `server_list`
    - if the server was unhealthy and there was a successful healthcheck
      - add the server back to the `server_list`

## Backend Server Selection Options

- Backend server with a weight and a backup server
  - NGINX example:

```
http {
    upstream backend {
        server backend1.example.com weight=5;
        server backend2.example.com;
        server 192.0.0.1 backup;
    }

    server {
        location / {
            proxy_pass http://backend;
        }
    }
}
```
- Least connections
  - A request is sent to a server with the least number of active connections
- Least time
  - A request is sent to a server with the lowest average response time and the lowest number of active connections
  - Time can be:
    - Time to receive the response header
    - Time to receive full response body



## Backend Server Selection Options (Cont.)

- Limiting the Number of Connections
  - *Maximum number of connections per backend server*
  - *Number of connections in the queue*

```
upstream backend {
    server backend1.example.com max_conns=3;
    server backend2.example.com;
    queue 100 timeout=70;
}
```
- Hash (ip hash, generic hash)
  - *A server to which a request is sent is determined from the client IP address or an arbitrary value (string, request URL, etc.)*
- Server Slow-Start
  - *This prevents a recently recovered server from being overwhelmed*
  - *During server slow-start, connections may time out*
    - *This may cause the server to be marked as failed again.*

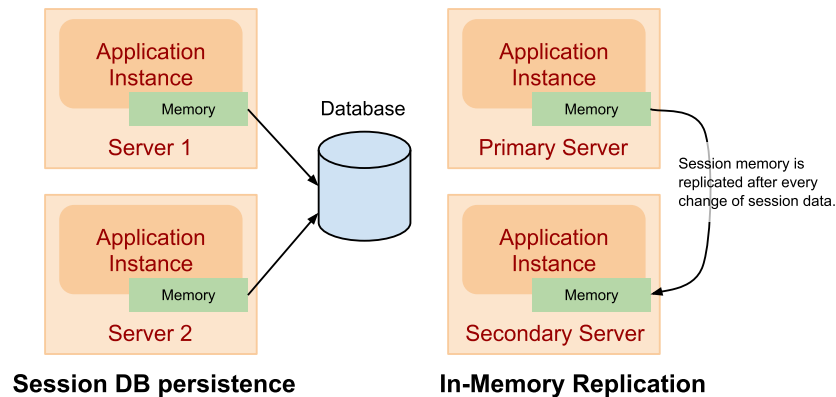
## Session Persistence

- Session Persistence
  - *Sticky cookie*
    - *A cookie defined by the load balancer for every client*

```
upstream backend {
    server backend1.example.com;
    server backend2.example.com;
    sticky cookie srv_id expires=1h domain=.example.com path=/;
}
```

  - *Sticky learn*
    - *LB finds a cookie by inspecting requests and responses*
    - *LB uses the cookie for subsequent redirection*

# Session State Persistence and Replication



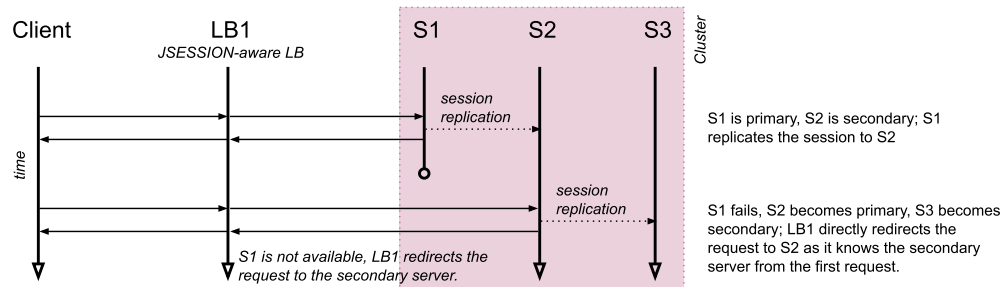
- Session DB persistence
  - Session information is maintained in the database
  - Does not require sticky sessions in LB
  - Implements `HttpSession` interface that writes data to the DB
- In-memory replication
  - A **primary server** holds a session state, the **secondary server** holds its replica.
  - Information about primary and secondary servers are part of `JSession`

## In-Memory Replication

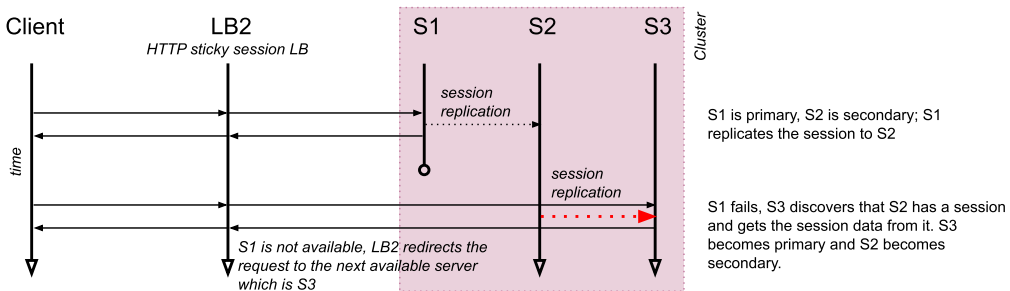
- Session format
  - It's a cookie
  - `JSESSIONID=SESSION_ID!PRIM_SERVER_ID!SEC_SERVER_ID!CREATION_TIME`
    - `SESSION_ID` – session id, generated by the server to identify memory associated with the session on the server
    - `PRIM_SERVER_ID` – ID of the managed server holding the session data
    - `SEC_SERVER_ID` – ID of the managed server holding the session replica
    - `CREATION_TIME` – time the session data was created/updated
- How LB uses this information
  - LB has information whether the server is running or not (via healthchecks)
  - if the primary server is running, it redirects the request there
  - if the primary server is not running, it redirects the request to the secondary server directly
  - if primary and secondary servers are not running, it redirect the request to any other server it has in the list – this may cause side effects!

# In-Memory Replication Scenarios

## Scenario A: JSession-aware load balancer



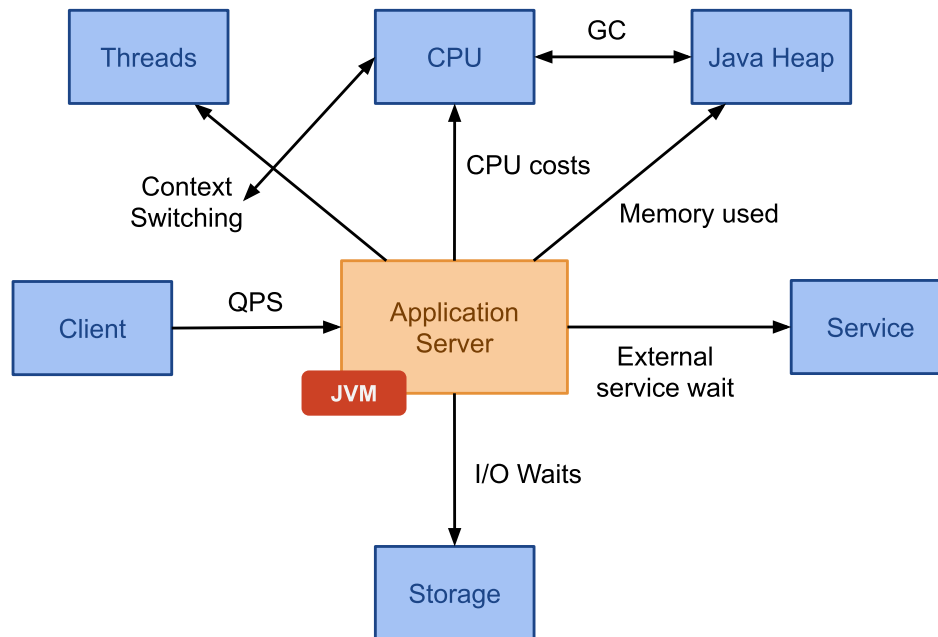
## Scenario B: HTTP sticky session load balancer



## Overview

- Definitions
- Load Balancers
- Performance Tuning

## Performance Limiting Factors



## Monitoring

- Important to understand performance
  - DevOps monitoring trends
- What you need
  - Collect → Filter → Store → View → **Tune**
  - Metrics, dashboards, alerting, log management, reporting, tracing capabilities
  - It is necessary to organize metrics well in order to understand what is going on
  - Start from a high-level process, detail to technical components
- Source
  - Application server
    - usually management beans with JMX interfaces
    - log files (access logs, server logs, etc.)
  - OS
    - many utilities available out of the box
    - open sockets, memory, context switches, I/O performance, CPU usage
  - Database
    - applications may write metrics to the DB
    - SQL scripts to collect metrics

## Monitoring Tools

- Commercial Monitoring Solutions
  - *Application server vendor usually offers a monitoring solution*
  - *AppDynamics, Oracle Enterprise Manager, Splunk*
  - *Google stackdriver, Amazon AWS CloudWatch*
- Open source examples
  - *Elasticsearch + LogStash + Kibana*
  - *InfluxDB + Telegraph + DataGraph*