

Middleware Architectures 1

Lecture 8: High Availability and Performance

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <https://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <https://vitvar.com/lectures>



Modified: Sun Dec 07 2025, 20:59:43
Humla v1.0

Good Performance

- What influences good performance?
 - *Number of users and concurrent connections*
 - *Number of messages and messages' sizes*
 - *Number of services*
 - *Infrastructure – capacity, availability, configuration, ...*
- How can we achieve good performance?
 - *Infrastructure*
 - *Scalability, failover, cluster architectures*
 - *Performance tuning*
 - *Application Server, JVM memory, OS-level tuning, Work managers configuration*
 - *Service configuration*
 - *Parallel processing, process optimization*

Overview

- **Definitions**
- Serving Requests
- Load Balancers
- HTTP in Kubernetes
- Performance Tuning

Definitions


- Scalability
 - *server scalability*
 - *ability of a system to scale – when input load changes*
 - *users should not feel a difference when more users access the same application at the same time*
 - **horizontal scaling**
 - *adding new instances of applications/servers*
 - **vertical scaling**
 - *adding new resources (CPU, memory) to a server instance*
 - *network traffic*
 - *bandwidth capacity influences performance too*
 - *service should limit the network traffic through caching*
- Availability
 - *probability that a service is operational at a particular time*
 - *e.g., 99.9987% availability – downtime ~44 seconds/year*
- SLA – Service Level Agreement
 - *Guarantee of service availability*
 - *When availability is below a guaranteed value, a customer can get a discount*

Definitions (Cont.)

- High Availability
 - *When a server instance fails, operation of the application can continue*
 - *Failures should affect application availability and performance as little as possible*
- Application Failover
 - *When an application component performing a job becomes unavailable, a copy of the failed object finishes the job.*
 - *Issues*
 - *A copy of the failed object must be available*
 - *A location and operational status of available objects must be available*
 - *A processing state must be replicated*
- Load Balancing
 - *Distribution of incoming requests across server instances*

RED Metrics: Rate

- Rate: Requests per Second (RPS/QPS)
 - *How many requests the service handles*
 - *A core server-side throughput metric*




- High Rate implications
 - *Caching can dramatically improve throughput*
 - *Even frequently changing data benefits from short-TTL cache*
 - *Load balancing and autoscaling rely on Rate metrics*

RED Metrics: Errors

- Errors: Failed Requests
 - *Rate of requests that fail or return unexpected result*
 - *Includes: 4xx/5xx codes, timeouts, retries, broken dependencies*
- Why errors matter
 - *Service health is mostly determined by error spikes*
 - *Errors propagate upstream causing cascading failures*
- Mitigation strategies
 - *Circuit breakers and retries*
 - *Graceful degradation*
 - *Better health checks & alerting*

RED Metrics: Duration

- Duration (Response Time)
 - *How long each request takes*
 - *Measured from the client or service side*




- High latency may indicate
 - *CPU-intensive service or inefficient algorithm*
 - *Bad configuration or resource limits*
 - *Slow downstream dependencies (DB, APIs)*
- Mitigation strategies
 - *Use asynchronous / background processing*
 - *Optimize hot paths, reduce blocking calls*

Overview


- Definitions
- **Serving Requests**
- Load Balancers
- HTTP in Kubernetes
- Performance Tuning

Blocking (Synchronous) I/O




- Inbound connection
 - A server creates a thread for every inbound connection
 - For example, 1K connections = 1K threads, big overhead
 - A thread is reserved for the entire duration of the request processing
- Outbound connection
 - A thread is blocked when outbound connection is made
 - When outbound connection is slow, the scalability is poor

Non-Blocking (Asynchronous) I/O



- Inbound connections
 - *The connection is maintained by the OS, not the server app*
 - *The Web app registers events, OS triggers events when they occur*
 - *The app may create working threads and controls their number*
- Outbound connections
 - *The app registers a callback that is called when the data is available*
 - *Event loop*

Components



- **Muxer** – component that handles communication via network sockets.
- **Request queue** – queue of requests to be processed.
- **Self-tuning thread pool** – a pool of threads in various states.
- **Work manager** – a configuration of **maximum threads** and a **capacity** that can be used to handle requests for a specific application/service.

Overview

- Definitions
- Serving Requests
- **Load Balancers**
- HTTP in Kubernetes
- Performance Tuning


Load Balancing

- Distributes a load to multiple app/object instances
 - *App instances run on different machines*
 - *Load sharing: equal or with preferences*
 - *Health checks*
- Types
 - *DNS-based load balancer*
 - *DNS Round Robin*
 - *NAT-based load balancer (Layer-4)*
 - ***Reverse-proxy load balancer (Layer-7)***
 - *application layer*
 - *Sticky sessions*
 - *JSession, JSession-aware load balancer*
 - *Client-side load balancer*
 - *LB run by a client*
 - *a client uses a replica-aware stub of the object from the server*

DNS-based Load Balancer


- DNS Round Robin
 - *A DNS record has multiple assigned IP addresses*
 - *DNS system delivers different IP addresses from the list*
 - *Example DNS A Record:*
`company.com A 147.32.100.71 147.32.100.72 147.32.100.73`
- Advantages
 - *Very simple, easy to implement*
- Disadvantages
 - *IP address in cache, could take hours to re-assign*
 - *No information about servers' loads and health*

Forward Proxy




- Forward Proxy
 - Centralized access control based on content
 - The client knows about the site it wants to access
 - Performs request on behalf of the client
 - Caches content to increase performance, limits network traffic
 - Filters requests or controls access based on destinations or origins
 - Widely used in private networks in companies
 - Most of the proxy servers today are Web proxy servers

Forward Proxy Sequence Diagram




- Forward Proxy
 - *There are 2 TCP sockets for every connection (client-FP, FP-server)*
 - *Client is not directly connected to the server*
 - *FP tunnels requests to the server*

Reverse Proxy




- Reverse Proxy
 - *Aggregates multiple request-response interactions with back-end systems*
 - *Processes the request on behalf of the client*
 - *The client does not know about the back-end systems*
 - *May provide additional capabilities*
 - *Data transformations*
 - *Security – authentication, authorization*
 - *Orchestration of communication with back-end systems*

Reverse Proxy Load Balancer



HTTP Sticky Sessions Example



- How to identify a server that hosts the session state
 - *Passive cookie persistence* – LB uses a cookie from the server
 - *Active cookie persistence* – LB adds its own cookie

Types of Load Balancers

- Software

- *Apache mod_proxy_balancer, NGINX*
 - *HTTP Session persistence – sticky sessions*
 - *Various configuration options*
- *WebLogic proxy plug-in*

```
1 <Location /soa-infra>
2     SetHandler weblogic-handler
3     WebLogicCluster czfmwapp03-vf:8001,czfmwapp04-vf:8001,czfmwapp05-vf:8001
4 </Location>
5
```

/soa-infra is a first part of an URL path that rules in this Location will be applied (this is a standard Apache configuration mechanism)

WebLogicCluster specifies the list of backend servers for load balancing

- Hardware

- *Cisco, Avaya, Barracude*

Round-Robin Algorithm

- Uses
 - request** – *client request with or without a cookie information*
 - server_list** – *a list of backend servers that can process the request*
 - rbinx** – *round robin index*
 - sticky_sessions** – *associative array of pairs <session_id,server>*
 - unhealthy_treshhold** – *a number of negative consecutive health checks before moving the server to the "unhealthy" state.*
- Round Robin Algorithm
 - *if **session_id** exist in the **request** and in **sticky_sessions***
 - *send the **request** to the server **sticky_sessions[session_id]***
 - *otherwise*
 - *send the **request** to the **rbinx** server in the **server_list***
 - *extract **session_id** from the response from the server*
 - *if the **session_id** exist, add a pair <**session_id**;server_list[rbinx]> to **sticky_sessions***
 - *increase **rbinx** by one or reset it to 0 if it exceeds the length of **server_list***

Health Check

- Health Check
 - For each server in the `server_list`
 - call the server's healthcheck endpoint
 - if a number of failed health checks for the server exceeds the `unhealthy_threshold`
 - remove the server from the `server_list`
 - if the server was unhealthy and there was a successful healthcheck
 - add the server back to the `server_list`

Backend Server Selection Options

- Backend server with a weight and a backup server

- *NGINX example:*

```
http {  
    upstream backend {  
        server backend1.example.com weight=5;  
        server backend2.example.com;  
        server 192.0.0.1 backup;  
    }  
  
    server {  
        location / {  
            proxy_pass http://backend;  
        }  
    }  
}
```

- Least connections

- *A request is sent to a server with the least number of active connections*

- Least time

- *A request is sent to a server with the lowest average response time and the lowest number of active connections*

- *Time can be:*

- *Time to receive the response header*

- *Time to receive full response body*

Backend Server Selection Options (Cont.)

- Limiting the Number of Connections

- *Maximum number of connections per backend server*
- *Number of connections in the queue*

```
upstream backend {  
    server backend1.example.com max_conns=3;  
    server backend2.example.com;  
    queue 100 timeout=70;  
}
```

- Hash (ip hash, generic hash)

- *A server to which a request is sent is determined from the client IP address or an arbitrary value (string, request URL, etc.)*

- Server Slow-Start

- *This prevents a recently recovered server from being overwhelmed*
- *During server slow-start, connections may time out*
 - *This may cause the server to be marked as failed again.*

Session Persistence

- Session Persistence

- *Sticky cookie*

- *A cookie defined by the load balancer for every client*


```
upstream backend {  
    server backend1.example.com;  
    server backend2.example.com;  
    sticky cookie srv_id expires=1h domain=.example.com path=/  
}
```

- *Sticky learn*

- *LB finds a cookie by inspecting requests and responses*

- *LB uses the cookie for subsequent redirection*

Session State Persistence and Replication



- Session DB persistence
 - Session information is maintained in the database
 - Does not require sticky sessions in LB
 - Implements `HttpSession` interface that writes data to the DB
- In-memory replication
 - A **primary server** holds a session state, the **secondary server** holds its replica.
 - Information about primary and secondary servers are part of `JSession`

Overview

- Definitions
- Serving Requests
- Load Balancers
- HTTP in Kubernetes
- Performance Tuning

HTTP in Kubernetes

- Kubernetes does not provide built-in Layer 7 (HTTP) routing
- Services expose Pods at TCP/UDP level
 - **ClusterIP**: *internal-only*
 - **NodePort**: *accessible via NodeIP:NodePort*
 - **LoadBalancer**: *cloud-provider load balancer*
- For advanced routing (URLs, hostnames, TLS)
 - *Kubernetes uses the Ingress API*
- Ingress works at the HTTP layer (L7) and provides:
 - *Path-based routing*
 - *Host-based routing*
 - *TLS termination*
 - *Centralized access point*

Ingress Resource

- An Ingress is a set of HTTP routing rules
- It defines how external HTTP/S traffic reaches cluster Services
- An Ingress needs an Ingress Controller
 - *NGINX Ingress Controller*
 - *HAProxy Ingress*
 - *Traefik*
 - *Contour*
- Controller watches Ingress objects and configures
 - *reverse proxy*
 - *load balancer*
 - *dynamic routing tables*

How Ingress Works

- Client sends HTTP request to public IP
 - *LoadBalancer or NodePort*
- Ingress Controller receives the request
 - *Runs as a Deployment inside the cluster*
 - *Listens on ports 80/443*
 - *Uses reverse proxy (NGINX/HAProxy/etc.)*
- Ingress rules decide where to send the request
 - *Host matching* → **example.com**
 - *Path matching* → **/api, /app**
- The controller forwards the request to the Service
 - *The Service then forwards to Pods*
- Ingress Controller handles
 - *TLS termination*
 - *HTTP routing*
 - *Load balancing across Pods*

Ingress Example

- Example Ingress routing traffic to the **busybox-tcp-service**


```
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: example-ingress
5    annotations:
6      nginx.ingress.kubernetes.io/rewrite-target: /
7  spec:
8    rules:
9      - host: busybox.example.com
10      http:
11        paths:
12          - path: /
13            pathType: Prefix
14            backend:
15              service:
16                name: busybox-tcp-service
17                port:
18                  number: 8080
```

- Requests to **http://busybox.example.com/** are routed to the Service
- Ingress Controller performs HTTP reverse proxying and load balancing across Pods

Overview

- Definitions
- Serving Requests
- Load Balancers
- HTTP in Kubernetes
- Performance Tuning

Performance Limiting Factors



Monitoring

- Important to understand performance
 - *DevOps monitoring trends*
- What you need
 - *Collect → Filter → Store → View → **Tune***
 - *Metrics, dashboards, alerting, log management, reporting, tracing capabilities*
 - *It is necessary to organize metrics well in order to understand what is going on*
 - *Start from a high-level process, detail to technical components*
- Source
 - *Application server*
 - *usually management beans with JMX interfaces*
 - *log files (access logs, server logs, etc.)*
 - *OS*
 - *many utilities available out of the box*
 - *open sockets, memory, context switches, I/O performance, CPU usage*
 - *Database*
 - *applications may write metrics to the DB*
 - *SQL scripts to collect metrics*

Monitoring Tools

- Commercial Monitoring Solutions
 - *Application server vendor usually offers a monitoring solution*
 - *AppDynamics, Oracle Enterprise Manager, Splunk*
 - *Google stackdriver, Amazon AWS CloudWatch*
- Open source examples
 - *Elasticsearch + LogStash + Kibana*
 - *InfluxDB + Telegraph + DataGraph*