


Middleware Architectures 1

Lecture 8: High Availability and Performance

doc. Ing. Tomáš Vitvar, Ph.D.

tomas@vitvar.com • @TomasVitvar • <https://vitvar.com>



Czech Technical University in Prague

Faculty of Information Technologies • Software and Web Engineering • <https://vitvar.com/lectures>



Modified: Mon Dec 09 2024, 07:02:53
Humla v1.0

Good Performance

- What influences good performance?
 - *Number of users and concurrent connections*
 - *Number of messages and messages' sizes*
 - *Number of services*
 - *Infrastructure – capacity, availability, configuration, ...*
- How can we achieve good performance?
 - *Infrastructure*
 - *Scalability, failover, cluster architectures*
 - *Performance tuning*
 - *Application Server, JVM memory, OS-level tuning, Work managers configuration*
 - *Service configuration*
 - *Parallel processing, process optimization*

Overview

- **Definitions**
 - *Serving Requests*
 - *Load Balancers*
 - *Performance Tuning*

Definitions


- **Scalability**
 - *server scalability*
 - *ability of a system to scale – when input load changes*
 - *users should not feel a difference when more users access the same application at the same time*
 - **horizontal scaling**
 - *adding new instances of applications/servers*
 - **vertical scaling**
 - *adding new resources (CPU, memory) to a server instance*
 - *network traffic*
 - *bandwidth capacity influences performance too*
 - *service should limit the network traffic through caching*
- **Availability**
 - *probability that a service is operational at a particular time*
 - *e.g., 99.9987% availability – downtime ~44 seconds/year*
- **SLA – Service Level Agreement**
 - *Guarantee of service availability*
 - *When availability is below a guaranteed value, a customer can get a discount*

Definitions (Cont.)

- High Availability
 - When a server instance fails, operation of the application can continue
 - Failures should affect application availability and performance as little as possible
- Application Failover
 - When an application component performing a job becomes unavailable, a copy of the failed object finishes the job.
 - Issues
 - A copy of the failed object must be available
 - A location and operational status of available objects must be available
 - A processing state must be replicated
- Load Balancing
 - Distribution of incoming requests across server instances

RED Metrics: Rate

- Rate: Requests per Second (RPS/QPS)
 - How many requests the service handles
 - A core server-side throughput metric




- High Rate implications
 - Caching can dramatically improve throughput
 - Even frequently changing data benefits from short-TTL cache
 - Load balancing and autoscaling rely on Rate metrics

RED Metrics: Errors

- Errors: Failed Requests
 - Rate of requests that fail or return unexpected result
 - Includes: 4xx/5xx codes, timeouts, retries, broken dependencies
- Why errors matter
 - Service health is mostly determined by error spikes
 - Errors propagate upstream causing cascading failures
- Mitigation strategies
 - Circuit breakers and retries
 - Graceful degradation
 - Better health checks & alerting

RED Metrics: Duration

- Duration (Response Time)
 - How long each request takes
 - Measured from the client or service side




- High latency may indicate
 - CPU-intensive service or inefficient algorithm
 - Bad configuration or resource limits
 - Slow downstream dependencies (DB, APIs)
- Mitigation strategies
 - Use asynchronous / background processing
 - Optimize hot paths, reduce blocking calls

Overview


- Definitions
 - *Serving Requests*
 - *Load Balancers*
 - *Performance Tuning*

Blocking (Synchronous) I/O




- Inbound connection
 - A server creates a thread for every inbound connection
 - For example, 1K connections = 1K threads, big overhead
 - A thread is reserved for the entire duration of the request processing
- Outbound connection
 - A thread is blocked when outbound connection is made
 - When outbound connection is slow, the scalability is poor

Non-Blocking (Asynchronous) I/O



- Inbound connections
 - The connection is maintained by the OS, not the server app
 - The Web app registers events, OS triggers events when they occur
 - The app may create working threads and controls their number
- Outbound connections
 - The app registers a callback that is called when the data is available
 - Event loop

Components



- **Muxer** – component that handles communication via network sockets.
- **Request queue** – queue of requests to be processed.
- **Self-tuning thread pool** – a pool of threads in various states.
- **Work manager** – a configuration of **maximum threads** and a **capacity** that can be used to handle requests for a specific application/service.

Overview

- Definitions
 - *Serving Requests*
 - *Load Balancers*
 - *Performance Tuning*


Load Balancing

- Distributes a load to multiple app/object instances
 - *App instances run on different machines*
 - *Load sharing: equal or with preferences*
 - *Health checks*
- Types
 - *DNS-based load balancer*
 - *DNS Round Robin*
 - *NAT-based load balancer (Layer-4)*
 - ***Reverse-proxy load balancer (Layer-7)***
 - *application layer*
 - *Sticky sessions*
 - *JSession, JSession-aware load balancer*
 - *Client-side load balancer*
 - *LB run by a client*
 - *a client uses a replica-aware stub of the object from the server*

DNS-based Load Balancer


- DNS Round Robin
 - A DNS record has multiple assigned IP addresses
 - DNS system delivers different IP addresses from the list
 - Example DNS A Record:
`company.com A 147.32.100.71 147.32.100.72 147.32.100.73`
- Advantages
 - Very simple, easy to implement
- Disadvantages
 - IP address in cache, could take hours to re-assign
 - No information about servers' loads and health

Forward Proxy




- Forward Proxy
 - Centralized access control based on content
 - The client knows about the site it wants to access
 - Performs request on behalf of the client
 - Caches content to increase performance, limits network traffic
 - Filters requests or controls access based on destinations or origins
 - Widely used in private networks in companies
 - Most of the proxy servers today are Web proxy servers

Forward Prox Sequence Diagram




- Forward Proxy
 - There are 2 TCP sockets for every connection (client-FP, FP-server)
 - Client is not directly connected to the server
 - FP tunnels requests to the server

Reverse Proxy




- Reverse Proxy
 - Aggregates multiple request-response interactions with back-end systems
 - Processes the request on behalf of the client
 - The client does not know about the back-end systems
 - May provide additional capabilities
 - Data transformations
 - Security – authentication, authorization
 - Orchestration of communication with back-end systems

Reverse Proxy Load Balancer



HTTP Sticky Sessions Example



- How to identify a server that hosts the session state
 - *Passive cookie persistence* – LB uses a cookie from the server
 - *Active cookie persistence* – LB adds its own cookie

Types of Load Balancers

- Software

- Apache `mod_proxy_balancer`, `NGINX`
 - HTTP Session persistence – sticky sessions
 - Various configuration options
- WebLogic proxy plug-in

```
1 <Location /soa-infra>
2     SetHandler weblogic-handler
3     WebLogicCluster czfmwapp03-vf:8001,czfmwapp04-vf:8001,czfmwapp05-vf:800
4 </Location>
5
```

`/soa-infra` is a first part of an URL path that rules in this `Location` will be applied (this is a standard Apache configuration mechanism)
`WebLogicCluster` specifies the list of backend servers for load balancing

- Hardware

- Cisco, Avaya, Barracude

Round-Robin Algorithm

- Uses

`request` – client request with or without a cookie information
`server_list` – a list of backend servers that can process the request
`rbinx` – round robin index
`sticky_sessions` – associative array of pairs `<session_id,server>`
`unhealthy_treshold` – a number of negative consecutive health checks before moving the server to the "unhealthy" state.

- Round Robin Algorithm

- if `session_id` exist in the `request` and in `sticky_sessions`
 - send the `request` to the server `sticky_sessions[session_id]`
- otherwise
 - send the `request` to the `rbinx` server in the `server_list`
 - extract `session_id` from the response from the server
 - if the `session_id` exist, add a pair `<session_id;server_list[rbinx]>` to `sticky_sessions`
 - increase `rbinx` by one or reset it to 0 if it exceeds the length of `server_list`

Health Check

- Health Check
 - For each server in the `server_list`
 - call the server's healthcheck endpoint
 - if a number of failed health checks for the server exceeds the `unhealthy_threshold`
 - remove the server from the `server_list`
 - if the server was unhealthy and there was a successful healthcheck
 - add the server back to the `server_list`

Backend Server Selection Options

- Backend server with a weight and a backup server
 - NGINX example:

```
http {
    upstream backend {
        server backend1.example.com weight=5;
        server backend2.example.com;
        server 192.0.0.1 backup;
    }

    server {
        location / {
            proxy_pass http://backend;
        }
    }
}
```
- Least connections
 - A request is sent to a server with the least number of active connections
- Least time
 - A request is sent to a server with the lowest average response time and the lowest number of active connections
 - Time can be:
 - Time to receive the response header
 - Time to receive full response body

Backend Server Selection Options (Cont.)

- Limiting the Number of Connections

- *Maximum number of connections per backend server*
- *Number of connections in the queue*

```
upstream backend {  
    server backend1.example.com max_conns=3;  
    server backend2.example.com;  
    queue 100 timeout=70;  
}
```

- Hash (ip hash, generic hash)

- *A server to which a request is sent is determined from the client IP address or an arbitrary value (string, request URL, etc.)*

- Server Slow-Start

- *This prevents a recently recovered server from being overwhelmed*
- *During server slow-start, connections may time out*
 - *This may cause the server to be marked as failed again.*

Session Persistence


- Session Persistence

- *Sticky cookie*
 - *A cookie defined by the load balancer for every client*

```
upstream backend {  
    server backend1.example.com;  
    server backend2.example.com;  
    sticky cookie srv_id expires=1h domain=.example.com path=/;  
}
```

- *Sticky learn*
 - *LB finds a cookie by inspecting requests and responses*
 - *LB uses the cookie for subsequent redirection*

Session State Persistence and Replication




- Session DB persistence
 - Session information is maintained in the database
 - Does not require sticky sessions in LB
 - Implements **HttpSession** interface that writes data to the DB
- In-memory replication
 - A **primary server** holds a session state, the **secondary server** holds its replica.
 - Information about primary and secondary servers are part of **JSession**

In-Memory Replication

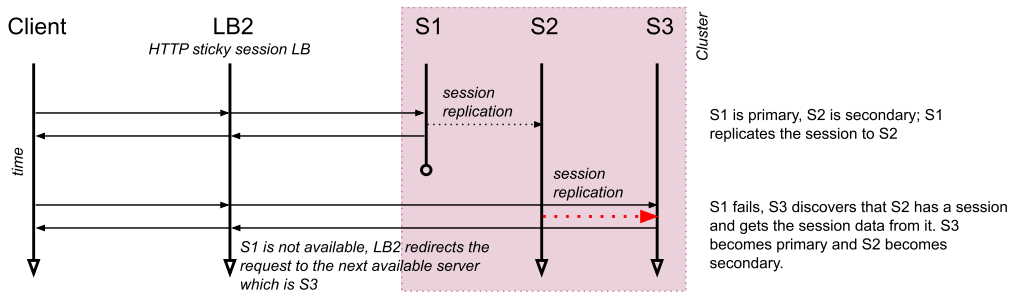
- Session format
 - It's a cookie
 - **JSESSIONID=SESSION_ID!PRIM_SERVER_ID!SEC_SERVER_ID!CREATION_TIME**
 - SESSION_ID** – session id, generated by the server to identify memory associated with the session on the server
 - PRIM_SERVER_ID** – ID of the managed server holding the session data
 - SEC_SERVER_ID** – ID of the managed server holding the session replica
 - CREATION_TIME** – time the session data was created/updated
- How LB uses this information
 - LB has information whether the server is running or not (via healthchecks)
 - if the primary server is running, it redirects the request there
 - if the primary server is not running, it redirects the request to the secondary server directly
 - if primary and secondary servers are not running, it redirect the request to any other server it has in the list – this may cause side effects!

In-Memory Replication Scenarios

Scenario A: JSession-aware load balancer




Scenario B: HTTP sticky session load balancer



Overview

- Definitions
 - *Serving Requests*
 - *Load Balancers*
 - *Performance Tuning*

Performance Limiting Factors



Monitoring

- Important to understand performance
 - DevOps monitoring trends
- What you need
 - Collect → Filter → Store → View → **Tune**
 - Metrics, dashboards, alerting, log management, reporting, tracing capabilities
 - It is necessary to organize metrics well in order to understand what is going on
 - Start from a high-level process, detail to technical components
- Source
 - Application server
 - usually management beans with JMX interfaces
 - log files (access logs, server logs, etc.)
 - OS
 - many utilities available out of the box
 - open sockets, memory, context switches, I/O performance, CPU usage
 - Database
 - applications may write metrics to the DB
 - SQL scripts to collect metrics

Monitoring Tools

- Commercial Monitoring Solutions
 - *Application server vendor usually offers a monitoring solution*
 - *AppDynamics, Oracle Enterprise Manager, Splunk*
 - *Google stackdriver, Amazon AWS CloudWatch*
- Open source examples
 - *Elasticsearch + LogStash + Kibana*
 - *InfluxDB + Telegraph + DataGraph*